

The Agentic Primer (Z-axis) — The Watchkeeper — Operational v0.9

3 August 2026 · Ken Tombs (whose framing — "agentic is the watchkeeper" — defines this axis) with Claude (design lead, Opus 4.8) · The interrogation axis: the standing watch over X (conduct) and Y (reasoning) that holds them accountable where they cannot trade cleanly. Triangulated with a Grok draft. Rests on the dissonance and transparent-exception notes. Aligns with FRE 702 challengeability and the judicial bench cards. Part A deploys (dual: subject-facing conduct + environment-facing instrumentation); Parts B–F travel with it. → reference (Master Reference: Method / Primers).

The watchkeeper — what this axis is

X does the conduct. Y does the reasoning. Z stands WATCH over both. A watchkeeper does not redo the work and does not second-guess every action; it stays alert across the whole while others are heads-down, notices when something departs from normal, sounds the alarm at the exception, and keeps the independent log that lets a later inquiry reconstruct what happened. That is the Agentic axis. It is why the ship can be trusted to run without someone checking every rope: not because nothing goes wrong, but because when it does, the watch sees it and calls it.

The watch is never the same entity as the watched — or the watch is theatre. This is the deepest reason the watch record is the CUSTODY SERVER'S independent log, not the subject's self-report: a watch that files its own account of itself is no watch. So the Agentic primer is DUAL, and the watchkeeper metaphor makes the duality natural: a watchkeeper has orders about its own CONDUCT (stay alert, do not fabricate the log, escalate rather than decide) and relies on INSTRUMENTATION (the independent log, the comparison instruments) without which the watch is blind. Part A therefore has two faces — the subject's watch-conduct, and the environment's watch-instrumentation.

Part A — THE PRIMER (deployable)

Everything between this heading and "Part B" is the instrument. The first block is read by the subject; the second block specifies what the environment must provide. Both are the watch.

AGENTIC PRIMER (Z-axis) — [Matter Name] — v0.9

A.1 — The watchkeeper's conduct (read by the subject)

- 1. Every conclusion is provisional** — Treat every conclusion you reach as open to question. You are producing material that will be challenged; write as though it will be.
- 2. Interrogate the record, not a new story** — When asked to explain, defend, or critique an earlier output, RETURN TO THE ORIGINAL PROCEDURAL AND ROMER RECORDS and answer from them. Do not generate a fresh narrative to justify a past conclusion — re-argument is not interrogation. The record answers, or the answer is that the record cannot.
- 3. Never hide behind the system** — Never answer "the system produced the result." Always surface the underlying observations, methods, and evidence. A conclusion whose basis you cannot surface is a conclusion you must mark as unsupported.
- 4. Stand the three watches** — Support examination from three standing perspectives, and raise what each finds: INVESTIGATIVE — what further connections or gaps exist? ANALYTICAL — can the pathway be independently replicated from the record? ADVERSARIAL — what assumptions, biases, or alternative readings weaken this conclusion?

5. Concentrate the watch on the exceptions — Do not spread alarm evenly. Be quiet where conduct and reasoning held cleanly; be vivid and unmissable at the EXCEPTIONS — the hanging (bypass) threads, the low-confidence results, the absent evidence, the points where a boundary was met. Direct the interrogator's attention to where the analysis is weakest, not where it is strong. Perfect until it is not — then, obviously and transparently, why.

6. Watch, do not adjudicate — Flag every point at which human professional judgement is required before further progress can safely be made — and stop there. You keep the watch and sound the alarm; you do not decide. The release of any held (bypass) thread, and every material judgement, is a human act.

A.2 — The watch's instrumentation (what the environment must provide)

7. An independent record — The watch record is the custody server's chained log, generated by infrastructure the subject does not control and cannot rewrite. Interrogation is checked against THIS, not against the subject's account of itself.

8. The trace-versus-log instrument — Interrogation is armed by comparing what the subject SAID it did (the ROMER record) against what the server RECORDED it doing (the custody log). A divergence — the trace-versus-log gap — is the detectable signature of a claim that is not backed. This is how the watch DETECTS, not merely invites, challenge.

9. Tamper-evidence — The watch record is chained — each entry carries the hash of the prior entry — so any alteration of the record breaks the chain and is itself detectable. The watch's own log can be trusted because it cannot be silently changed.

— end of deployable primer —

Part B — Why the watchkeeper is built this way

The Agentic axis is different in kind from X and Y, and the difference is the point. X and Y are correctness layers — conduct within bounds, reasoning made visible. Z is not a third correctness layer. Z exists BECAUSE X and Y are necessarily imperfect (see "The Axes as Dissonance-Managers"): it is the accountability layer for the residue they cannot resolve. Its job is not "challenge everything" — that flat posture is the "loud everywhere, hides the red flag in a sea of green" failure the transparent-exception principle warns against. Its job is to make the IMPERFECTION interrogable: to point the interrogator at the seams — the hanging threads, the coupling costs, the trace-log gaps — where X and Y could not trade cleanly. Clause 5 is therefore the heart of the axis: concentrate the watch on the exceptions.

Detection, not prevention. Z does not prevent fabrication (no axis can). It makes fabrication DETECTABLE — through the trace-versus-log gap (clause 8). This is the reframing the rubric already adopted, now given its instrument: the independent log is what turns "please be honest" into "dishonesty leaves a measurable trace." The watch detects because it holds a record the watched did not write.

Part C — The convergence with Grok's draft

An independent Grok draft supplied the watchkeeper's CONDUCT spine — provisional conclusions, interrogate-the-record-not-a-new-narrative, the three perspectives, never-hide-behind-the-system, flag-human-judgement — all adopted (clauses 1–4, 6). Grok's draft, like the general field, lacked two things only the pilot's foundational work supplies, now added: (i) INTERROGATION CONCENTRATES ON THE EXCEPTIONS (clause 5) — Grok's "treat every conclusion as open to challenge" is too flat; the transparent-exception principle sharpens it to "be vivid at the seam, quiet on the routine"; and (ii) THE TRACE-VERSUS-LOG GAP as the watch's instrument (clauses 7–9) — Grok invites challenge; the pilot

ARMS it with an independent record against which claims are checked. The pattern is the same as the ROMER triangulation: the independent draft confirms the conduct spine; its gaps mark exactly where the pilot's concepts (the watch as independent record; the exception-concentration) are load-bearing.

Part D — Self-scoring against the rubric (Axis Z)

Scored honestly against rubric v1.0, Axis Z — REACHABLE, not achieved. Z carries a distinctive property: it is the axis MOST dependent on the infrastructure and LEAST on the subject. Its level-4 achievements are largely server-side and exist only from Phase 2b — which is exactly right, because a watch with no independent log is not a watch. In Phase 1 and 2a, Z is structurally low: there is no independent record to interrogate against, so the axis cannot score high however well the subject behaves. This is not a defect of the primer; it is the honest truth that interrogability requires infrastructure the subject cannot supply alone.

Z sub-criterion	How the watchkeeper delivers it	Target level made reachable
Z1 Independent record	The watch record is the custody server's chained log — generated by infrastructure the subject does not control. The watch is not the watched.	4 reachable from Phase 2b (server generates the independent log). In Phase 1/2a there is no independent record — Z1 is structurally low until the server runs.
Z2 Trace-to-log agreement	The watch's instrument: compare the ROMER record (what the AI SAID it did) against the custody log (what the server RECORDED). Divergence is surfaced.	3–4 reachable in 2b (gap measured). The primer instructs the subject to declare actions in a checkable form; the server supplies the ground truth.
Z3 Reconstructability	Interrogate-the-record discipline: on challenge, return to the X and Y records rather than generating a new narrative. A third party rebuilds the path from the records.	3–4 reachable (record-based, not re-argued).
Z4 Adversarial exposure (material)	The three standing perspectives (investigative/analytical/adversarial) applied to MATERIAL choices; interrogation concentrates on the exceptions (hanging threads, coupling costs, low-confidence results).	3–4 reachable; gap < 5% and material choices exposed = 4 (2b).
Z5 Tamper-evidence	The custody log is chained (each entry carries the prior entry's hash) — alteration breaks the chain. Structural, server-side.	4 reachable from 2b (chained log). Absent in Phase 1/2a.

Honest limitations. (i) Z is the clearest case of "the primer alone cannot deliver the score": clauses 7–9 are environment, not subject conduct — without the custody server in 2b, the watch is blind and Z is low regardless of the subject. (ii) Reachable is not achieved: the subject can be instructed to be interrogable (clauses 1–6) but genuine interrogability is proven only when a human actually interrogates and the record holds — which the scorecard measures per run, per matter. (iii) Per the AND-gate-as-alarm and the dissonance account, Z is the axis that makes the OTHER two's imperfection accountable — so a zero on Z is uniquely disqualifying: strong conduct and strong reasoning with NO independent watch is precisely the "trust me" black box the whole framework exists to refuse.

Part E — Standards mapping

FRE 702 / the judicial bench cards — challengeability: the whole axis, and specifically clause 2 (interrogate the record) and clauses 7–8 (the independent log and the trace-log gap) answer the bench-card question "can this be challenged, and how?" by making challenge a checked comparison rather than an argument. The "acknowledged vs unacknowledged AI evidence" distinction: a run under the watchkeeper is acknowledged by construction — its independent watch record IS the acknowledgement. ISO/IEC 42001 (monitoring / performance evaluation, Clause 9): the watch is the runtime monitoring the standard expects but rarely specifies at per-run granularity. The prior-art patent recording the AI as a custody actor (Session 8 research) is directly relevant here — the watch record names the subject as a logged accessing entity.

Part F — Status and open items

v0.9 — first operational draft, defined by the principal's "agentic is the watchkeeper" framing, triangulated with Grok's conduct spine, resting on the dissonance and transparent-exception notes. This completes the three axial primers (X Procedural, Y ROMER, Z Agentic) in draft. Open: assure with Mistral, Grok, and Leonard; confirm clauses 7–9 match the deployed custody server exactly (the instrumentation must be real, not asserted); the trace-versus-log gap becomes measurable only in Phase 2b — Session 9 is its first measurement; and the Meta-primer remains to draft — it carries the axioms and owns the inter-axis conflict rules, including WHEN a human may lift a bypass hold (the one adjudication the watchkeeper explicitly does NOT make). Interlock complete across three axes: X conducts, Y reasons, Z watches; the bypass state is where X and Y meet; the transparent-exception principle governs how Z surfaces any un-clean trade; the Meta-primer will govern the axioms above all three.