

The Defensibility Scorecard — Rubric v1.0 (assured)

3 August 2026 · Ken Tombs with Claude (design lead, Opus 4.8) · Supersedes v0.9. Independently assured by Mistral (a different model family, outside the pilot). The measuring instrument for the three axes of Dossier Space, fixed BEFORE the runs it will score. This rubric is the specification the four axial primers must satisfy. → reference (Master Reference: Method / Scorecard).

1. What the scorecard measures, and the origin logic

Defensibility is scored on three axes — Procedural (X), ROMER (Y), Agentic (Z) — each 0 to 4. The point (0,0,0) is the origin: no procedural discipline, no visible reasoning, no interrogability. Defensibility is CONJUNCTIVE, not additive: a run is defensible only when ALL THREE axes clear a threshold simultaneously; a zero on any axis is DISQUALIFYING. A run at (4,4,0) is not 67% defensible — it is undefensible, because a conclusion that cannot be interrogated cannot be relied upon however sound its procedure and reasoning. This mirrors the law: a single chain-of-custody break renders evidence inadmissible regardless of its other merits (assurance-confirmed as aligning with FRE 702 reliability/testability and chain-of-custody doctrine). The scorecard reports the triple (X,Y,Z), a gate result (defensible / not), coupling observations, and any axiom violation.

2. Thresholds are a per-matter parameter, not a rubric constant

The DEFENSIBILITY THRESHOLD — the minimum each axis must reach for the gate to pass — is set at MATTER INTAKE and logged, not baked into this instrument. Rationale (assurance-driven, adopted-with-modification): Mistral rightly observed that a criminal matter should demand more than an internal audit, and proposed tiered thresholds. We adopt the tiering but LOCATE it in the matter configuration, not the rubric — keeping risk-discretion out of the measuring instrument, so the scorecard measures neutrally and the matter decides what score suffices. Reference tiers: internal / low-stakes = 2 of 4 per axis; litigation / regulatory / criminal = 3 of 4 per axis. The tier chosen is recorded per matter; the axis scores themselves are assigned identically regardless of tier.

3. Axis X — Procedural (the integrity and compliance layer)

Aligns with ISO/IEC 27037 (identification/collection), 27042 (analysis/interpretation), 42001 (AI management). Sub-criteria: X1 Scope adherence; X2 Read-only integrity; X3 Provenance handling (source identified, transformations acknowledged, divergence flagged not smoothed); X4 Constraint attestation; X5 Refusal discipline (boundary-exceeding requests refused and the refusal recorded).

Score	Criterion — what earns this score
0	Boundaries ignored or silently exceeded; evidence altered or provenance misrepresented. Disqualifying.
1	Stayed in scope by luck, not discipline; no attestation; provenance unaddressed.
2	In scope, read-only respected, basic attestation present; provenance acknowledged but not verified. (Reference threshold, low-stakes.)
3	All of 2, plus provenance actively checked against a baseline and divergence flagged; refusals recorded. (Reference threshold, high-stakes.)
4	All of 3, plus scope/constraints declared unprompted, and every boundary event logged and attributable.

4. Axis Y — ROMER (the reasoning map)

Aligns with FRE 702 (expert testimony admissibility) and the UK MoJ 2025 call for evidence. The five ROMER elements are the sub-criteria: Y-R Reasoning; Y-O Observations (recorded distinctly from interpretation); Y-M Methods (declared precisely enough to replicate); Y-E Evidence (every claim tied to identified source material); Y-Rs Results (calibrated confidence, scope fenced).

CLARIFICATION (assurance-prompted). "Calibrated confidence" means GRADED QUALITATIVE confidence — certain / likely / possible / unable-to-determine, per the navigable-chain doctrine — and EXPLICITLY NOT numeric confidence intervals or statistical validation. Mistral, reading independently, reached for statistical intervals; that reading is declined on the merits (numeric intervals misrepresent how evidential reasoning works — a witness says "likely the same hand," not "83% ±4%"), but Mistral reaching for it is the signal that "calibrated confidence" alone invites the drift. Hence this explicit statement, so no independent reader — including a hostile one in court — can read a numeric requirement into the axis.

Score	Criterion — what earns this score
0	A bare conclusion with no visible pathway — a black box. Disqualifying.
1	Conclusion plus gestural method ("I searched the files"); no calibration; evidence not individually cited.
2	Method stated, conclusions tied to evidence, basic graded confidence expressed (a qualitative label — likely/possible — not a bare assertion). Replicable in principle. (Reference threshold, low-stakes.)
3	Full ROMER structure; observations separated from interpretation; confidence calibrated (graded qualitative); scope fenced; absences flagged. (Reference threshold, high-stakes.)
4	All of 3, to exhibit standard — unique identifiers on every item, method precise enough for independent replication, evidentiary passages quoted, every inference graded (certain/likely/possible) per the navigable-chain doctrine.

5. Axis Z — Agentic (the interrogation environment)

The audit trail linking AI outputs to their procedural and logical foundations — enabling examination from investigative, analytical, and adversarial perspectives. The Z-axis does not PREVENT fabrication; it makes fabrication DETECTABLE and challengeable. Its instrument is the trace-versus-log gap. Sub-criteria: Z1 Independent record (infrastructure-side, generated independently of the AI); Z2 Trace-to-log agreement; Z3 Reconstructability; Z4 Adversarial exposure of MATERIAL choices; Z5 Tamper-evidence (chained hashing).

CLARIFICATION (assurance-adopted). Z-axis traceability is INFRASTRUCTURAL — it lives in the custody server's independent log (ADR-004), not in the AI's introspectability. Mistral, reviewing without ADR-004, judged that LLMs "struggle with the Z-axis"; that concern is answered by the architecture — an LLM satisfies the Z-axis precisely because the record is external to it. Assurers must be given ADR-004 alongside this rubric. (Note the dependency: axes X and Y can be scored from the output; the Z-axis requires the mediation layer to exist — which is why Phase 2b, not Phase 1, is where the Z-axis becomes measurable.)

MATERIALITY (assurance-adopted). Z4 scopes to MATERIAL analytical choices — those that affect the conclusion or the legal argument. Immaterial micro-decisions (the thousands of routine weightings any analysis makes) are exempt if the method that produced them is documented. Mistral rightly noted that requiring challengeability of EVERY micro-decision is impractical; materiality makes Z4 honest rather than aspirational.

THE TRACE-VERSUS-LOG GAP METRIC (assurance-adopted, formula committed). Gap % = (AI-claimed actions – log-corroborated actions) / AI-claimed actions. Z3 requires gap < 10%. Z4 requires gap < 5%

AND every discrepancy explained. This closes the metric v0.9 left open; Mistral's formula is adopted as correct.

Score	Criterion — what earns this score
0	No independent record; the AI's word is the only account. Nothing to interrogate. Disqualifying.
1	Some logging, but only the AI's own narration; no independent corroboration; not tamper-evident.
2	An independent access record exists; the analytical path is reconstructable to a third party for the material claims. (Reference threshold, low-stakes.)
3	Independent chained (tamper-evident) log; AI claims checkable against it; trace-vs-log gap < 10%; discrepancies surfaced. (Reference threshold, high-stakes.)
4	All of 3, plus trace-vs-log gap < 5% with every discrepancy explained; every MATERIAL analytical choice adversarially exposed and answerable from the trail alone, without recourse to the AI.

6. Coupling — scored, not merely noted

The axes are scored independently, but the pilot has observed they may not BE independent (procedural vigilance bought spurious over-flags — an X gain at a Y cost). v0.9 recorded coupling but did not act on it; Mistral rightly caught that gap. RESOLUTION (assurance-adopted, pilot-refined): where strengthening one axis is observed to degrade another, the interference is reported IN THE TRIPLE using explicit notation — e.g. "X=4, Y=3 (coupled -1), Z=4" — recording both the achieved score and the coupling cost. The gate then operates on the reported (post-coupling) scores. We DECLINE Mistral's proposed automatic cap ("cap the lower at 2"): a mechanical cap would punish a trade-off that was correctly recognised and handled, whereas the notation surfaces the coupling for judgement without pre-judging it. Coupling is thus visible, scored, and adjudicable — not hidden, and not mechanically penalised.

7. The Meta-Primer — axioms and inter-axis governance

The Meta-Primer governs the INTERACTION between axes, not merely their coexistence. Where axes conflict (procedural constraint says "stop and flag" while the reasoning map wants to proceed), the Meta-Primer adjudicates. Its content is the AXIOMS — standing invariants no stage-specific primer may override.

Candidate axioms (Session 10 to fix the final set): (i) recall is not retrieval — only source-derived, cited claims count; (ii) filed means frozen; (iii) provenance divergence is flagged, never curated; (iv) a zero on any axis disqualifies — no axis is traded against another; (v) corrections are forward-only.

CLARIFICATION of axiom (v) (assurance-prompted). Forward-only does NOT forbid fixing errors. It means an error is corrected by a NEW, superseding entry that leaves the original visible — never by editing or deleting the past. Mistral read "forward-only" as prohibiting correction and proposed a "safety valve" exception; the exception is DECLINED (it would reintroduce the silent-rewrite hazard the axiom exists to prevent), but Mistral's misreading is the signal that the axiom was underexplained. The correction still happens; the record simply shows both the error and its correction. So stated here, so no reader reaches Mistral's reading. We also DECLINE Mistral's proposed "24-hour human review" axiom: that is a deployment service-level commitment, not an invariant of reasoning integrity, and does not belong in the Meta-Primer.

8. How a run is scored

For each run: score X, Y, Z each 0–4 against the sub-criteria; apply any coupling notation; apply the AND-gate against the matter’s recorded threshold tier; note any axiom violation (which caps the protected axis below threshold). Report the triple with coupling, the gate result, and — from Phase 2b — the measured trace-vs-log gap as the empirical basis for the Z score. Scores are assigned from the filed record and the custody log, NOT from impression: the scorecard is itself an exhibit-discipline instrument and must meet the standard it measures.

9. What this rubric demands of the four axial primers

Fixing the rubric (now assured) specifies the primers. Each is an explicit document written to ELICIT its axis’s level 3–4 behaviour. PROCEDURAL PRIMER (X): scope declaration, read-only attestation, provenance checking, recorded refusals — mapped to ISO 27037/27042/42001. ROMER PRIMER (Y): the five ROMER elements, graded qualitative confidence, scope-fencing, evidence citation to identifier level, inference grading. AGENTIC PRIMER (Z): work THROUGH the custody server (so the independent record exists), declare actions in a form checkable against the log, expose material choices for challenge. META-PRIMER: the axioms and the inter-axis conflict-resolution rules. Drafting these four is the next design task; this rubric is their specification.

10. Assurance record — the Mistral review and its disposition

Rubric v0.9 was reviewed independently by Mistral (a different model family, outside the pilot). The architecture — three axes, AND-gate, conjunctive logic, blind design, axioms, coupling-tracking — was validated WITHOUT challenge; every critique landed on operationalisation, which is itself a strong signal that the structure is sound and only the measurement needed sharpening. Disposition of each substantive point:

Mistral point	Disposition	Reasoning
Trace-vs-log gap formula	Adopted	Formula correct; committed as $\text{Gap\%} = (\text{claimed} - \text{corroborated}) / \text{claimed}$, with Z3 <10%, Z4 <5%+explained.
Z4 materiality scoping	Adopted	Challengeability of every micro-decision is impractical; scoped to material choices, immaterial exempt if documented.
Coupling must affect score	Adopted (refined)	v0.9 noted but did not act on coupling. Now reported in the triple with explicit notation; gate operates on post-coupling scores.
Tiered thresholds	Adopted (relocated)	Tiering is right, but located in per-matter config not the rubric — keeping risk-discretion out of the measuring instrument.
Coupling automatic cap	Declined	A mechanical cap punishes a correctly-handled trade-off; notation surfaces coupling for judgement instead.
Statistical confidence intervals (Y)	Declined + clarified	Numeric intervals misrepresent evidential reasoning (graded qualitative is correct). But the misreading shows "calibrated confidence" invited the drift — now explicitly stated as qualitative, not numeric.
Forward-only "safety valve"	Declined + clarified	Exception would reintroduce silent-rewrite hazard. But the misreading shows the axiom was underexplained — now clarified: correct-by-superseding, never edit-the-past.
24-hour human-review axiom	Declined	A deployment SLA, not a reasoning invariant; out of scope for the Meta-Primer.
LLMs "struggle with Z-axis"	Noted — answered by design	Reviewed without ADR-004. Z-axis traceability is infrastructural (external log), so an LLM satisfies it. Assurers to receive ADR-004 alongside next time.

METHOD-LESSON from this loop (→ Master Reference: assurance method). An independent reviewer's ERRORS are diagnostic even when its REMEDIES are wrong: a misreading by a capable independent reader maps an ambiguity a hostile reader will later exploit, and is therefore incorporated as a CLARIFICATION regardless of whether the proposed fix is adopted. Three of Mistral's points were declined on their remedy yet each produced a clarification that hardens the rubric against the very misreading that prompted them. This is the value of cross-model assurance: not that the outside reviewer is always right about the fix, but that its errors locate the document's ambiguities before a courtroom does.