

Dossier Space Pilot Note

What the Pilot Is Showing Against the Emerging Professional Position on AI-Assisted Disclosure

Status: Working observation

Date: August 2026

Purpose: To record emerging lessons from the Dossier Space pilot against current developments in criminal disclosure and AI-assisted evidence review.

1. Why this note

The Dossier Space pilot was not designed specifically to test the developing UK criminal disclosure reforms. Its purpose has been broader: to explore how people and AI might work together over substantial bodies of evidence while retaining custody, traceability, assurance and human judgement.

Recent professional and policy developments nevertheless provide an increasingly relevant external reference point.

Jonathan Fisher KC's *Disclosure in the Digital Age* examined the challenges created by the growth of digital evidence and produced 45 recommendations for modernising the disclosure regime. Sir Brian Leveson's *Independent Review of the Criminal Courts: Part 2* subsequently considered disclosure within the wider problem of criminal court efficiency. On 14 July 2026 the Government brought these strands together in *Modernising the Criminal Disclosure Regime*, its response to Fisher and to the disclosure chapter of Leveson's review.

Jake Eddison-Cook's practitioner analysis, *Paint by Numbers – Disclosure by Algorithm*, published by Peters & Peters on 27 July 2026 and subsequently reproduced by the Society for Computers & Law, identifies a particularly important practical issue: even where the final decision formally remains human, technology may increasingly determine the material on which that human decision is based.

That is very close to the problem being explored experimentally through Dossier Space.

2. What Dossier Space has encountered

The pilot has shown that using AI over evidential material is not simply a question of whether an AI model can find relevant information.

In practice, difficult questions have repeatedly concerned the surrounding evidential process:

- What corpus was actually searched?
- What search instruction or seed phrase was used?
- What did the search return?
- What may have been missed?
- Were results pooled or deduplicated?
- What metadata was retained?
- Did the search terminate normally?
- Was a timeout or technical failure visible?
- Were intermediate results preserved?
- Can the search later be reconstructed?
- Can another person challenge the search by changing its assumptions?

These issues became particularly apparent during the Session 9 investigation and its subsequent reconciliation work.

A relatively simple evidential question generated different candidate populations depending on the retrieval mechanism used. The earlier search produced one population; a later custody-based search identified additional candidates.

The difference did not in itself overturn the original evidential finding. It did, however, require explanation, reconciliation and preservation.

That experience appears significant.

The principal operational risk was not that AI produced an obviously absurd answer. The more subtle risk was that a technically plausible process could produce an incomplete evidential population without the incompleteness necessarily being apparent to the person receiving the result.

3. Where the pilot aligns positively with the emerging professional position

3.1 Human judgement remains important — but is not sufficient

The emerging professional position continues to emphasise human responsibility and oversight.

Dossier Space supports that principle, but the pilot suggests that “human in the loop” is not, by itself, an adequate assurance mechanism.

A human cannot exercise meaningful judgement over evidence that was never presented to them.

Assurance therefore has to extend upstream from the final decision into the process by which the evidential population is assembled.

In Dossier Space terms:

Evidence corpus → governed retrieval and analysis → inspectable evidential population → professional judgement

Human judgement remains at the end of the chain, but the chain leading to it must also be open to scrutiny.

3.2 Audit trails need to record more than the final AI answer

The pilot has progressively required preservation of:

- queries and parameters;
- result populations;
- source paths;
- custody information;
- hashes;
- processing steps;
- operator interventions;
- failures and timeouts;
- reconciliation decisions.

This is strongly consistent with the professional debate identified by Eddison-Cook around configuration, prompting, competent operation, accuracy, audit trails and scrutiny of AI-assisted disclosure.

The pilot suggests that an AI-generated report alone is a poor evidential record.

The process by which the report acquired its evidential population may be almost as important as the report itself.

3.3 Reproducibility is becoming a practical legal issue

One of the strongest pilot lessons has been the value of being able to reconstruct what happened.

Where two searches produce different populations, it should be possible to establish whether that difference arose from:

- search logic;
- corpus coverage;

- indexing;
- metadata;
- software behaviour;
- operator intervention;
- timeout or failure;
- or a genuine difference in the question asked.

This is more than conventional software logging.

In an evidential environment, reproducibility contributes to the ability to explain why particular material was — or was not — placed before a decision-maker.

3.4 The pilot supports separating evidence-finding from evidential judgement

Dossier Space has increasingly moved towards an architecture in which AI assists in finding, organising, relating and testing evidence without being treated as the final legal decision-maker.

That distinction now appears increasingly important.

The AI may help establish the evidential field.

The professional determines its legal significance.

ROMER and the associated assurance process can preserve the reasoning and evidence path between those activities.

4. Where the pilot exposes concerns

4.1 Apparently successful AI processes can conceal incompleteness

A search can complete, produce convincing results and support a plausible narrative while still having missed relevant material.

Nothing necessarily announces that omission.

This may be a more difficult practical risk than obvious AI hallucination because the result may look entirely credible.

4.2 Tool choice can affect the evidential population

The reconciliation work showed that different retrieval mechanisms can return different candidate populations.

Statements such as “the AI searched the evidence” are therefore insufficiently precise.

Relevant questions include:

- which tool;
- which corpus;
- which configuration;
- which search instruction;
- and with what known limitations?

4.3 Scale creates both the case for AI and its assurance problem

Fisher's review arose substantially from the increasing volume and complexity of digital material in modern investigations. The Government has now accepted the need to modernise the regime around that reality.

AI-assisted review becomes attractive precisely because manual inspection at scale becomes increasingly impractical.

But the same scale also makes complete independent checking difficult.

Re-running the entire process manually defeats much of the purpose of using AI.

Assurance therefore has to be designed into the process rather than appended afterwards.

4.4 Technical failure becomes evidentially significant

During the pilot, timeouts, incomplete searches and infrastructure problems could not simply be treated as IT inconveniences.

If a failed or incomplete search affects what evidence is subsequently considered, the technical event may influence the evidential process.

It therefore needs to be recorded as part of that process.

5. An additional requirement emerging from the pilot: challengeability

Current discussions commonly refer to transparency, accountability, auditability and human oversight.

The Dossier Space pilot suggests another useful property:

challengeability.

Auditability asks:

Can we establish what the system did?

Reproducibility asks:

Can we perform substantially the same process again?

Challengeability asks:

Can another competent party test whether the process should have been performed differently?

That might involve:

- changing search terms;
- testing alternative assumptions;
- examining material outside an AI-selected population;
- comparing retrieval methods;
- inspecting exclusions;
- or repeating a search independently.

This is particularly significant in an adversarial legal environment.

Eddison-Cook identifies essentially this problem from the defence perspective: if a defence team believes an AI-assisted prosecution process has missed something, what can it meaningfully ask to inspect, rerun or challenge?

Dossier Space does not presently answer the legal or procedural question of what access must be granted.

It does suggest what technical and evidential artefacts would need to exist for such a challenge to be meaningful.

6. Metadata and the evidential field

Fisher's review contains a dedicated metadata annex, and the Government's response accepts the wider role of metadata schedules in modernising disclosure.

This appears particularly relevant to Dossier Space.

The pilot has already shown that the evidential value of an item cannot always be separated from information about:

- where it came from;
- how it was found;
- when it was retrieved;
- its relationship to other records;
- and how it moved through the analysis process.

Metadata is therefore not merely administrative information around the evidence.

In an AI-assisted environment, it can become part of the evidence needed to establish how an evidential population was formed.

7. Potential relevance to smaller defence practices

Eddison-Cook raises a significant concern that the move towards sophisticated AI-assisted disclosure may widen the capability gap between prosecution bodies and smaller defence practices.

The Dossier Space pilot may point towards one possible response.

It does not necessarily follow that a smaller practice needs an equivalent copy of the prosecution's technology platform.

Instead, it may need sufficient capability to:

1. understand how an evidential population was created;
2. inspect records describing that process;
3. formulate alternative searches or challenges;
4. run proportionate independent tests;
5. preserve those results; and
6. present them to a legal professional for judgement.

This is consistent with the pilot's emerging use of relatively simple local tools, structured tasks, custody controls and AI guidance rather than dependence upon a single opaque end-to-end platform.

This remains a proposition to be tested rather than a demonstrated capability.

8. What the pilot does **not** establish

Care is needed not to overstate these observations.

The Dossier Space pilot does not demonstrate that:

- current prosecution disclosure systems are unreliable;
- the Government's proposed arrangements are inadequate;
- Dossier Space complies with CPIA disclosure requirements;
- AI-assisted disclosure can be assured solely through technical logging;
- every difference between retrieval mechanisms is legally material;
- or the current pilot architecture is ready for live operational deployment.

The Enron corpus is a research dataset, not a criminal case.

The pilot has deliberately exposed weaknesses in architecture, custody and process so they can be identified and understood.

Its present value is primarily diagnostic and experimental.

9. Positive and negative implications for the profession

Seen against the current professional position, the Dossier Space experience produces a mixed but constructive picture.

Positive

AI can make large evidential populations more tractable.

Search and analysis can be repeated quickly.

Alternative hypotheses can be explored.

Large numbers of records can be organised into human-manageable evidential structures.

Custody and provenance can, in principle, be built into the workflow.

Machine activity can be logged more extensively than many conventional manual review processes.

A well-designed AI-assisted environment may therefore increase rather than reduce evidential visibility.

Negative

Those benefits are not automatic.

Searches that appear successful may still be incomplete.

Different tools can produce different evidential populations.

Failures can remain hidden unless explicitly instrumented.

An apparently clear AI answer can obscure a complicated or imperfect retrieval process.

At large scale it may become impossible for the professional making the final judgement to inspect personally more than a small proportion of the underlying material.

The profession therefore faces a shift in assurance responsibility: from checking individual documents towards also understanding and assuring the systems by which those documents were selected.

10. Emerging conclusion

Recent professional and policy developments appear to strengthen rather than weaken the rationale for Dossier Space.

The disclosure regime is moving towards a position in which technology, including AI, may increasingly participate in determining what material reaches a human decision-maker.

The Dossier Space pilot suggests that governance applied only at the point of final professional judgement is therefore too late.

Attention also needs to be given to the evidential journey:

what was available → what was searched → how it was searched → what was found → what was not found → what changed → what failed → what was presented → what was judged.

The positive finding from the pilot is that these stages can be made substantially more observable and reconstructible.

The negative finding is that this does not happen automatically. Even relatively straightforward pilot searches have exposed differences in retrieval, metadata, custody and infrastructure that could easily disappear behind an apparently convincing AI-generated answer.

The emerging Dossier Space proposition can therefore be stated simply:

When AI influences what evidence a person sees, the evidential process should remain inspectable, reconstructible and challengeable.

That does not replace professional judgement.

It seeks to ensure that professional judgement is exercised over an evidential field whose origins can still be examined.

References and professional context

Eddison-Cook, Jake (2026). *Paint by Numbers – Disclosure by Algorithm*. Peters & Peters, 27 July 2026.

Primary practitioner commentary considered in this note. Subsequently reproduced by the Society for Computers & Law.

Fisher KC, Jonathan (2025/2026). *Disclosure in the Digital Age: Independent Review of Disclosure and Fraud Offences*. Home Office. Originally published 20 March 2025; GOV.UK publication updated 14 July 2026. The review contains 45 recommendations and includes dedicated annexes on governance and metadata.

HM Government (2026). *Modernising the Criminal Disclosure Regime*. Home Office, 14 July 2026. Government response to Fisher’s *Disclosure in the Digital Age* and to Chapter Five, Disclosure, of Leveson’s *Independent Review of the Criminal Courts: Part 2, Volume One*.

Leveson, Sir Brian (2026). *Independent Review of the Criminal Courts: Part 2*. Ministry of Justice, 4 February 2026.

HM Government (2026). *AI to speed up justice under major disclosure reforms*. GOV.UK, 14 July 2026. Policy announcement accompanying the disclosure reforms.

Dossier Space Pilot (2026). Sessions 9–10 and associated custody, search, reconciliation and assurance records.

Status: Working pilot observation. This note records experimental findings and their possible relevance to emerging professional practice. It is not legal advice and does not claim operational compliance with the criminal disclosure regime.