

DOSSIER SPACE

Model Workstation Build — Overview Record

Session record for Design Lead (DL) / Operator records

Date	28 August 2026
Environment	Remote sandbox (isolated from operator workstation)
Status	INITIAL_SETUP complete — ready for small test corpus
Corpus status	No live corpus present (by design). Structure is corpus-agnostic.
Source brief	Dossier Space Workstation Build Induction Brief for DL (28 Aug 2026)

1. Purpose of this record

This document records what was actually built in response to the Design Lead induction brief. It serves as a concise, reconstructible summary for the operator’s records. It does not itself constitute a formal analytical session or ROMER.

2. What was requested

The induction brief instructed DL to design and establish a sovereign, method-neutral Dossier Space capable of working directly over a complete evidential corpus (initially the full raw Enron set, ~120 GB), without pre-committing to RAG, vector databases or any other retrieval fashion. The design had to be portable, reconstructible, and usable by a competent reviewer or successor AI without reliance on the originating conversation.

Immediate deliverables requested (§13 of the brief):

- 1. Proposed physical and logical folder tree
- 2. Root governing files and starter Dossier Space Primer
- 3. Session-template folder implementing expected-versus-actual capture
- 4. Simple logging design
- 5. Custody baseline approach suitable for ~120 GB
- 6. Method-library structure (direct, RAG, Diplomats, AHP/ANP, etc.)
- 7. Portability / bootstrap note
- 8. Short validation plan for the first direct-corpus Enron test

3. What was delivered

All eight requested deliverables were completed inside the isolated sandbox. The resulting model is self-describing, uses relative paths, and contains no live large corpus.

3.1 Top-level structure created

```
DOSSIER_SPACE_ROOT/  
├─ 00_CONTROL/      Environment status, portability, logging design, validation plan  
├─ 01_RAW_CORPUS/    Immutable source only (placeholder README)  
├─ 02_CUSTODY/       Acquisition record + tiered custody baseline approach  
├─ 03_PRIMERS_METHODS/ Starter Primer
```

— 04_SESSIONS/	Session template with expected / observed / ROMER
— 05_WORKING/	Disposable scratch area
— 06_LOGS/	Central logging location
— 07_EVIDENCE/	Preserved result sets
— 08_FINDINGS/	Finding register location
— 09_ROMER/	ROMER artefacts
— 10_OUTPUTS/	Downstream reports & exhibits
— 11_ASSURANCE/	Independent review records
— 12_TOOLS/	Optional helper scripts
— 13_METHOD_LIBRARY/	Nine method cards
— 99_ARCHIVE/	Superseded material

3.2 Root governing files

- **README.md** — orientation for any new human or AI
- **DIRECTIVE.md** — controlling directive and non-negotiable principles
- **DOSSIER.yaml** — identity, paths, roles, configuration
- **ROLES.md** — role register and authority boundaries
- **CHANGELOG.md** — material changes to structure/controls
- **KNOWN_LIMITATIONS.md** — technical, evidential and operational limits
- **SESSION_REGISTER.csv** — formal session tracking
- **FINDINGS_REGISTER.csv** — stable finding IDs and status
- **METHOD_REGISTER.md** — available methods and approval state

3.3 Session template (expected-versus-actual)

Located at 04_SESSIONS/_TEMPLATE_DS-YYYYMMDD-SNN/. Contains:

- plan/EXPECTED.md — completed before execution
- execution/OBSERVED.md — completed after execution; reconciles against plan
- ROMER.md — produced from the captured record, not from memory
- results/, findings/, assurance/, logs/ subfolders
- Session README describing the minimal workflow

3.4 Method library

Nine method cards under 13_METHOD_LIBRARY/, each stating when suitable, minimum inputs, main risks, expected outputs and assurance requirements:

- Direct / native filesystem & database interrogation (default)
- Conventional full-text / regex / metadata search
- RAG / embeddings / semantic retrieval (selective only)
- Diplomatics & documentary authenticity/provenance
- AHP / ANP and structured decision methods
- Timeline & sequence reconstruction
- Entity / relationship / network analysis
- Statistical & anomaly analysis
- Classification / clustering (selective)

3.5 Supporting control documents

- **00_CONTROL/LOGGING_DESIGN.md** — lightweight, reconstructible logging principles
- **02_CUSTODY/CUSTODY_BASELINE_APPROACH.md** — tiered strategy practical at ~120 GB scale
- **00_CONTROL/PORTABILITY.md** — how to recreate the model on another workstation

- **00_CONTROL/VALIDATION_PLAN_ENRON.md** — first-validation steps for the environment itself

4. Design principles observed

- **Immutable source** — raw corpus remains untouched; work occurs around it
- **Direct access is the default substrate** — RAG and other techniques are selectable, never mandatory
- **Expected vs actual** — built into every formal session template
- **Method visibility & failure visibility** — parameters, paths, errors and interventions are first-class
- **Role separation** — Human Operator, Design/Leader AI, Execution tools, Dossier Space AI, Assurance
- **Portability** — relative paths, plain-text controls, no vendor lock-in
- **Reconstructibility** — a competent reviewer can understand the environment without the original chat

5. Explicit limitations of this build session

- The full raw Enron corpus (~70–120 GB) was never present and could not be loaded into this chat environment.
- No absolute paths or machine-specific drive letters were hard-coded; those remain for the target SSD.
- No live analytical session was executed against real evidence.
- No access exists (or was attempted) to the operator's local workstation, SSD or network.
- All artefacts exist only inside the isolated remote sandbox used for this conversation.

6. Capability confirmation (what this session demonstrated)

This session confirmed that a capable AI operating under the induction brief can:

- Read and interpret a detailed method-neutral design brief
- Produce a complete, numbered, human-intelligible folder tree
- Create the full set of governing files, primers, registers and method cards
- Implement the expected-versus-actual capture model inside a reusable session template
- Document logging, custody, portability and first-validation approaches
- Remain strictly within the sandbox boundary and correctly report the absence of access to the operator's machine
- Leave a coherent, reconstructible record of what was requested and what was actually delivered

7. Recommended next steps

9. Copy or transfer the entire DOSSIER_SPACE_ROOT tree onto the target SSD / workstation.
10. Place a small test corpus under 01_RAW_CORPUS/ (preserve original structure and filenames).
11. Complete 02_CUSTODY/ACQUISITION.md and a lightweight custody baseline.
12. Create the first formal session from the template (e.g. DS-20260828-S01).
13. Load the Starter Primer and execute a bounded direct-access task.
14. Complete EXPECTED → OBSERVED → ROMER and preserve result artefacts.
15. Only after the environment itself is validated, consider scale tests or selective supplementary methods.

Working principle

*Evidence remains native. Methods are selectable. Actions are visible.
Material artefacts are captured. ROMER reconciles intent with execution. Assurance tests the record.*