

Dossier Space Workstation Induction and Design Brief

For Design Lead (DL)

Purpose: create a reusable sovereign Dossier Space on the new SSD and validate it against the full raw Enron corpus (~120 GB).

Status: Working Directive / Model Environment Design Brief | 28 August 2026

DIRECTIVE STATEMENT

Create a sovereign, method-neutral Dossier Space in which an AI can work directly over a complete evidential corpus, invoke additional analytical techniques only where justified, and leave a coherent, reconstructible record of what it was expected to do, what it actually did, what evidence it used, what methods it applied, what changed or failed, and what conclusions can properly be supported.

1. Purpose of this induction

This brief instructs DL to design and establish the next Dossier Space working environment on the new SSD. The immediate test corpus is the full raw Enron dataset, expected to be approximately 120 GB. The design must not be Enron-specific. It should be capable of being reproduced on another suitable workstation for another evidential corpus without redesigning the underlying method.

The objective is no longer merely to prove that an AI can search a large corpus. Earlier pilot sessions have already shown that useful analysis is possible. The next objective is to make the whole working environment coherent, visible, governable and reconstructible.

2. Design proposition

Dossier Space should assume that capable sovereign AI platforms can increasingly interrogate native files, directories, databases and machine tools directly. Chunking, embeddings, RAG, indexing and other retrieval techniques remain available, but they are not to be treated as mandatory architecture.

Direct access is the default substrate. Technique follows evidential need.

- Do not pre-commit the environment to RAG, vector databases, chunking or any other retrieval fashion.
- Do not prevent those techniques being used when they materially improve retrieval, recall, speed or repeatability.
- Treat Diplomats, AHP/ANP, statistical analysis, graph/network analysis, timeline reconstruction, entity tracing, conventional search and similar approaches as selectable methods.
- Record why a method was selected, how it was applied, what inputs it used, what outputs it created and what effect it had on the finding.

- Keep the source corpus unchanged; analytical work occurs around it, not inside it.

3. Lessons to carry forward from the pilot

The earlier Dossier Space sessions indicate that the principal risks are increasingly methodological and evidential rather than simply computational. The workstation design should therefore respond directly to those lessons.

Pilot lesson	Design consequence
Immutable source	The raw corpus must remain untouched and separately identifiable from all working artefacts.
Method visibility	Search terms, scripts, parameters, paths, exclusions, time windows and transformations must be preserved.
Expected vs actual	The system should capture the intended method before execution and reconcile it with what actually occurred afterwards.
Result preservation	Significant result sets should be persisted, identifiable and hashed where appropriate; do not rely on transient console output.
Failure visibility	Timeouts, partial searches, path errors, access failures, retries and operator interventions are evidential events and should be logged.
Reconciliation	Counts and conclusions must be capable of being reconciled across logs, result sets, findings and reports.
Role separation	Human operator, leader/design AI, execution/coding tools and external assurance should remain distinguishable in the record.
Documentation restraint	Live documentation should record material facts without becoming so intrusive that it distorts the work being observed.
Reconstruction	A competent reviewer should be able to understand what happened without relying on the originating AI conversation or human memory.

4. Required top-level Dossier Space structure

DL should create a simple numbered top-level structure that remains intelligible to a human looking at the disk without specialist software. The names below are a recommended baseline; DL may improve naming where there is a clear reason, but the functional separation should be preserved.

```
DOSSIER_SPACE_ROOT/
├── 00_CONTROL /
├── 01_RAW_CORPUS/
├── 02_CUSTODY /
├── 03_PRIMERS_METHODS/
├── 04_SESSIONS/
├── 05_WORKING/
├── 06_LOGS/
├── 07_EVIDENCE/
├── 08_FINDINGS/
```

```
|— 09_ROMER/  
|— 10_OUTPUTS/  
|— 11_ASSURANCE/  
|— 12_TOOLS/  
|— 13_METHOD_LIBRARY/  
|— 99_ARCHIVE/
```

4.1 00_CONTROL

- Environment README and operating directive.
- Case/dossier identity and configuration.
- Role register and authority boundaries.
- Session numbering convention.
- Current status, open issues and known limitations.
- Change log for the Dossier Space structure itself.

4.2 01_RAW_CORPUS

- The complete source corpus only.
- Read-only wherever practical.
- No scripts, notes, output or temporary files inside the corpus tree.
- Original source directory structure and filenames preserved.

4.3 02_CUSTODY

- Acquisition/source record.
- Corpus manifests.
- File counts, sizes and directory inventories.
- Cryptographic hashes or a documented tiered hashing strategy appropriate to corpus scale.
- Integrity checks and later verification records.
- Any known source anomalies or limitations.

4.4 03_PRIMERS_METHODS

- Dossier-level Primer.
- Session/question Primers.
- CIPDA plans or equivalent approved method statements.
- ROMER templates.
- Search/search-family definitions.
- Approved taxonomies and analytical criteria.

4.5 04_SESSIONS

Each formal analytical session should have its own self-contained subfolder, e.g. DS-YYYYMMDD-SNN. A session folder should provide a local route to its plan, actual execution record, results, findings and assurance references without duplicating immutable evidence unnecessarily.

4.6 05_WORKING

- Scratch calculations, temporary extracts, staging files and intermediate transforms.
- Clearly disposable or non-evidential material.
- Subfolders or naming conventions that identify the creating session.
- No final evidence should exist only here.

4.7 06_LOGS

- Execution logs.
- Command/tool invocation logs.
- Search parameters and timestamps.
- Errors, timeouts and exceptions.
- Operator interventions and reruns.
- Environment/version information where material.

4.8 07_EVIDENCE

- Preserved result sets and extracts supporting findings.
- Copies or references to relevant source records.
- Hashes and manifests for preserved evidential artefacts.
- Evidence packages sufficient for later verification without polluting the raw corpus.

4.9 08_FINDINGS

- Candidate findings.
- Accepted findings.
- Rejected/superseded findings with reason.
- Stable finding IDs.
- Links/references to supporting evidence, methods, sessions and assurance status.

4.10 09_ROMER

- Session ROMERs.
- Finding-level ROMER where warranted.
- Dossier-level synthesis ROMER.
- Explicit reconciliation of planned method against actual execution.

4.11 10_OUTPUTS

- Reports, timelines, tables, exhibits, visualisations and deliverables.
- Outputs are downstream products, not substitutes for the evidence and method record.

4.12 11_ASSURANCE

- Independent review notes.
- Scorecards.
- Reconciliation checks.
- Acceptance/rejection records.
- Replications and challenge tests.

4.13 12_TOOLS

- Local scripts and utilities approved for use in the Dossier Space.
- Versioned where possible.
- Accompanied by minimal instructions, dependencies and tests.
- Tool output should flow to working/log/evidence locations, not into the raw corpus.

4.14 13_METHOD_LIBRARY

A method-neutral library of optional techniques. Each method should have a short method card explaining when it is suitable, minimum inputs, main risks, expected outputs and assurance requirements.

- Direct/native filesystem and database interrogation.
- Conventional full-text/regex/metadata search.
- RAG / embeddings / semantic retrieval.
- Diplomatics and documentary authenticity/provenance analysis.
- AHP/ANP and other structured decision methods.
- Timeline and sequence reconstruction.
- Entity/relationship/network analysis.
- Statistical and anomaly analysis.
- Classification/clustering where justified.

4.15 99_ARCHIVE

- Superseded but retained material.
- Closed session packages.
- Old tool versions or deprecated methods where retention is necessary.
- Nothing should be archived merely to hide unresolved discrepancies.

5. Minimum governing files

At the root, DL should establish a small set of human-readable files that make the Dossier Space self-describing. Plain text/Markdown/JSON/YAML are preferred where practical because they are portable and machine-readable.

File	Purpose
README.md	What this Dossier Space is, how it is organised, and the first instructions for a new human or AI.
DIRECTIVE.md	The controlling directive statement and non-negotiable operating principles.
DOSSIER.yaml	Dossier ID, title, corpus paths, status, roles, dates, root paths and configuration.
ROLES.md	Human operator, design/leader AI, coding/execution tools, Dossier Space AI role, assurance role and authority boundaries.
CHANGELOG.md	Material changes to environment, structure, methods or controls.
KNOWN_LIMITATIONS.md	Known technical, evidential and operational limitations.
SESSION_REGISTER.csv or .jsonl	Formal session ID, purpose, status, dates and principal outputs.
FINDINGS_REGISTER.csv or .jsonl	Stable finding ID, status, source session, evidence references and assurance state.
METHOD_REGISTER.md or .yaml	Available methods, versions and approved use state.

6. The Dossier Space Primer

DL should create a starter Primer that is loaded or presented to any AI before substantive work begins. It should not attempt to answer the evidential question itself. Its purpose is to establish the environment, method discipline and reporting expectations.

6.1 Starter Primer content

1. Identify the Dossier Space root and confirm the raw corpus location before analysis.
2. Treat the raw corpus as evidential source material; do not alter it.
3. Establish the current question, scope, date/time boundaries, entities and exclusions.
4. State what method is proposed and why it is suitable. Direct/native interrogation is the default; supplementary techniques may be invoked only where they add value.
5. Before execution, record what is expected to happen: intended searches/actions, permitted sources, expected outputs, success/failure conditions and known uncertainties.
6. During execution, preserve material commands, parameters, paths, result sets, failures, retries and deviations in the designated locations.
7. After execution, record what actually happened and reconcile it against the pre-execution expectation.
8. Distinguish observation from inference and inference from conclusion.
9. Preserve sufficient evidence to reproduce or challenge a material finding.
10. Complete ROMER using the actual working record rather than reconstructing it from memory after the event.
11. Do not hide contradictory results, failed searches, negative findings or unresolved discrepancies.
12. Report limitations proportionately and escalate when the evidence or method does not support a defensible conclusion.

7. Before/after capture model

This should be designed into the workspace, not left to operator memory. Every formal session should have a lightweight pre-execution and post-execution record.

Before: expected	After: observed	Feeds ROMER
Question, scope and constraints	Actual scope reached; exclusions or boundary changes	Reasoning / Methods
Planned sources and access paths	Sources actually accessed; failures or unavailable paths	Observations / Evidence
Planned searches, tools or methods	Commands, parameters, techniques and versions actually used	Methods
Expected result artefacts	Actual result sets, counts, hashes and locations	Evidence
Expected success and assurance checks	Checks actually completed; deviations and exceptions	Observations / Results
Known uncertainties	New uncertainties, contradictions and unresolved gaps	Reasoning / Results

8. ROMER role in the new environment

ROMER is to be informed by the captured record, not treated as a narrative written after the analytical work. Use the canonical formulation: Reasoning, Observations, Methods, Evidence, Results.

- Reasoning: why the approach and technique were chosen, including alternatives or constraints.
- Observations: what was encountered in the corpus and execution environment, including absences, anomalies and failures.
- Methods: what was intended, what was actually done, and any material variance between the two.
- Evidence: the source records, result sets, logs, manifests and preserved artefacts supporting the work.
- Results: what can properly be concluded, together with limits, unresolved issues and assurance state.

The planned-versus-actual record is therefore one of the principal inputs into ROMER and should be machine-readable where practical.

9. Portability and model Dossier Space requirements

DL should design for transferability from the outset. A second competent user should be able to create the same structure on another machine, point it at a different raw corpus, load the Dossier Primer, and understand how to begin without requiring access to the originating conversations.

- Use relative paths within the Dossier Space wherever possible.
- Avoid hard-coding the current drive letter except in a local machine configuration file.
- Keep dependencies explicit and minimal.
- Prefer open or plain-text control records.
- Provide bootstrap/setup scripts where they reduce manual error, but do not make the environment unintelligible without those scripts.
- Document tool and model dependencies separately from the evidential structure so AI platforms can be substituted later.
- Ensure the directory layout remains usable with ordinary operating-system tools.
- Allow future case-specific method modules without altering the core structure.

10. First validation: full raw Enron corpus

Once the SSD and Dossier Space structure are ready, the initial validation should use the full raw Enron corpus in situ. The aim is not immediately to repeat every prior pilot analysis. First validate the environment itself.

13. Confirm corpus completeness sufficiently for the test and record its physical location, size, file count and custody baseline.
14. Confirm read-only or equivalent protection for the raw corpus.
15. Confirm the AI can directly navigate and interrogate the corpus at machine level without first creating a RAG store or chunked duplicate.
16. Run a small known search from the earlier pilot to validate access and logging.
17. Confirm that expected-vs-actual records are generated and reconcile with logs and result artefacts.
18. Confirm result sets can be preserved and later re-opened independently.
19. Confirm a ROMER can be produced directly from the captured session record.
20. Have assurance review reconstruct the test from the Dossier Space artefacts rather than from chat history.
21. Only then test scale, performance and the selective use of supplementary methods such as RAG, Diplomats or AHP where an evidential question genuinely warrants them.

11. What DL should not optimise prematurely

- Do not build a vector database merely because the corpus is large.
- Do not chunk the entire corpus as a default preprocessing stage unless measurement demonstrates a need.
- Do not create a complex orchestration layer before direct machine-level work has been measured.
- Do not make the design dependent on one vendor, model or agent framework.
- Do not allow convenience logging to substitute for preserved evidential result sets.
- Do not bury the working method inside opaque automation that a reviewer cannot reconstruct.

12. Definition of a successful first model

The first model Dossier Space is successful if a fresh competent user or AI can enter the SSD, understand the governing directive and roles, identify the untouched source corpus, execute a bounded analytical task using an appropriate method, locate the full record of what was planned and what actually occurred, identify the evidence supporting the result, reconstruct any deviations or failures, and reproduce or challenge the finding without relying on undocumented memory or the original conversation.

Design test: If the AI technology were replaced tomorrow, would the Dossier Space still make sense, preserve the evidential record, and allow another capable system to continue the work? If yes, the architecture is probably at the right level.

13. Immediate deliverables requested from DL

22. A proposed physical and logical folder tree for the new SSD, including any local machine configuration layer.
23. The root governing files and starter Dossier Space Primer.
24. A session-template folder implementing the expected-versus-actual capture model.
25. A simple logging design that captures material AI/tool actions without making normal work unusable.
26. A custody baseline approach suitable for a ~120 GB raw corpus.
27. A method-library structure allowing direct access, RAG, Diplomats, AHP/ANP and future techniques to be invoked selectively.
28. A portability/bootstrap note explaining how the model Dossier Space could be recreated on another workstation.
29. A short validation plan for the first direct-corpus Enron test.

14. Working principle

Evidence remains native. Methods are selectable. Actions are visible. Material artefacts are captured. ROMER reconciles intent with execution. Assurance tests the record.