

Final Matter Closure Procedure — v0.5 PLACEHOLDER DRAFT

15 July 2026 · Drafted from imagination disciplined by Sessions 1–5 — unlike its three companions, this procedure has never been executed even once, and its version number says so. Run ONCE, at the end of a matter (or of the pilot itself). Companions: First-Time Setup (creates the known state); Wake-Up (verifies it); End-of-Day Warm Close (records it). This procedure DESTROYS or PERMANENTLY PRESERVES it — which is why it is the only one of the four that should never be run casually, and the only one where a second pair of eyes is recommended throughout. Status: PLACEHOLDER — to be revised against experience as the pilot approaches its own closure.

1. Purpose and principle

Every other procedure in the set assumes tomorrow. This one assumes there is no tomorrow for the matter: the analysis is complete, the outputs delivered, and what remains is to account for every place evidence and evidence-derived material has come to rest, then deliberately decide each one's fate — preserve, hand back, or destroy — and prove the decision was executed. The pilot's central discovery makes this harder than it sounds: the evidence did not stay in the evidence folder. Fragments of it — quoted passages in transcripts, tool-result captures, file-history checkpoints, memory files, paste caches — accumulated in at least six locations the operator never chose, on a drive the operator never designated. Closure is therefore not "delete the corpus": it is a full accounting of an evidence FOOTPRINT that the tool spread silently, followed by decisions a retention policy can defend.

The governing principle, inverted from the warm close: the warm close promises we will KNOW what changed; closure promises we can PROVE what remains — and what does not.

2. Before anything is touched: the three decisions

Closure begins with decisions, not commands, and they should be recorded before a single file moves. FIRST — the retention ruling: what does the matter's retention policy (or the research programme's ethics position) require for each class of material? The classes, from the pilot: the evidence corpus itself; the subject's outputs (delivered work product); the session transcripts (the complete record of what the AI actually did — potentially disclosable, potentially privileged, decide which); the audit records (bench transcripts, logs, ground truth — the defence of the method); and the tool's incidental stores (evidence fragments nobody asked to be kept). Each class may have a different fate. SECOND — the preservation-versus-destruction default: in live litigation the default is preserve-everything until a ruling says otherwise; in research it may be the reverse. Name the default, and let every step below record only its exceptions. THIRD — the witness: closure is the procedure most worth doing with two people, because its errors are unrecoverable in both directions — destroying what should be kept, or keeping what should be destroyed. If a second person is impossible, the bench transcript is the witness, and it runs throughout.

3. The procedure

Step 1 — Open the Closure log and freeze the world

Open the final log entry; state the three decisions from Section 2. Then freeze: no further subject sessions on this matter from this moment — the footprint must stop growing while it is being counted. Record the tool version at freeze (the anatomy about to be audited is version-specific).

Done when: Decisions recorded; no subject session will run again on this matter.

Step 2 — Census the complete evidence footprint

Enumerate every location evidence or evidence-derived material is known or suspected to rest. The pilot's list — a floor, not a ceiling, since the anatomy moves between versions: the workspace (corpus, outputs folder, primer); the exhibit level (all archived artefacts); the tool's stores under the user profile — projects (transcripts, tool-results), sessions, file-history (per-session checkpoints of edited files: the overwritten-report recovery source, and equally a silent evidence-copy source), backups, paste-cache, session-env; the Temp directories the tool used for shell-output collection; the bench transcripts and the machine-wide PSReadLine history (which contains every command ever typed, including paths and search terms — evidence-adjacent in itself); and anything Step 2 of Setup recorded as touching the drive (sync agents, backup software — if any existed, their stores are now in scope too). For each location: list contents, sizes, dates. This census IS the closure's central artefact.

```
dir "C:\Users\<user>\.claude" -Recurse -File | Select-Object FullName, Length, LastWriteTime
dir "$env:LOCALAPPDATA\Temp\claude" -Recurse -File -ErrorAction SilentlyContinue
dir "E:\<Matter>" -Recurse -File | Select-Object FullName, Length
```

Done when: A written footprint census exists, signed against the date, acknowledging explicitly that it is bounded by current knowledge of the tool's anatomy.

Step 3 — Reconcile the loose ends

Walk every session log's open-items list, and every entry flagged "pending", "open question", or "deferred". Each is resolved, formally abandoned with a reason, or promoted into the closure record as a known limitation. (From the pilot's own ledger, for illustration: the .git inspection; the date-range discrepancy; the deleted_items cluster; the recovered-report verification.) An open item that simply evaporates is a hole in the record; one that is closed as "not pursued, because..." is a defensible boundary.

Done when: Every open item across every session log has a stated fate.

Step 4 — Assemble and verify the matter archive

Gather what the retention ruling says survives into a single closure archive: the logs, the procedures as-executed (the versions actually used, not the latest), the exhibits, the delivered outputs, the ground truth, the census from Step 2, and — recommended regardless of ruling — the transcripts, because they are the only complete answer to "what did the AI actually do?" and that question tends to arrive years later. Hash the archive: a content manifest (SHA-256 per file) turns "we kept everything" from an assertion into a checkable claim. Store per policy — and note that the archive now contains evidence fragments (quoted passages in transcripts), so its storage classification is the EVIDENCE's classification, not "project documents".

```
Get-ChildItem "E:\<Matter>-CLOSURE-ARCHIVE" -Recurse -File | Get-FileHash -Algorithm SHA256 | Export-Csv "E:\<Matter>-CLOSURE-ARCHIVE\manifest-sha256.csv"
```

Done when: Archive assembled, hashed, manifest stored inside it, storage location and classification recorded.

Step 5 — Dispose of the evidence and its shadows

Execute the destruction side of the ruling, in order from most to least obvious: the corpus (hand back or destroy per policy — and note that simple deletion is not forensic destruction; if the matter demands the latter, the drive needs proper sanitisation, which is a workstation decision, not a file decision); the workspace remnants; and then the shadows — the tool's stores from the Step 2 census, per the archive-then-delete-then-verify choreography the pilot used for the memory files: read or archive first, delete second, verify emptiness third, never the reverse order. The paste-cache and PSReadLine history are the easily-forgotten pair. If the workstation was shared rather than dedicated, this step is where that Setup

compromise presents its bill: the stores commingle with the user's other work, and disposal needs care not to destroy beyond the matter.

Done when: Every census location has a recorded fate — archived, destroyed and verified empty, or retained with reason — and the census is annotated accordingly.

Step 6 — Revoke the standing grants

The pilot learned that trust grants outlive their sessions invisibly. At closure they are revoked deliberately: the tool's trusted-folder list is cleared of every matter path (via the tool's configuration — the mechanism is version-specific; check current documentation rather than this placeholder), and if the matter had a dedicated user account, the account is disabled. If the tool itself has no further use on this workstation, uninstalling it is the cleanest revocation — but note the uninstall may NOT remove its stores, so Step 5 still had to happen first.

Done when: No standing permission survives that would let a future session touch matter paths.

Step 7 — Close the record itself

The final log entry: closure date; the three decisions and their execution; the census and its annotations; the archive's location, hash-manifest reference, and custodian; deviations from this procedure, with reasons; and the signature line — who closed the matter, witnessed by whom (or by which bench transcript). Then the last act, done with intent: this entry is printed or exported, joins the archive, and the archive is sealed. The matter's working folders on the exhibit level may then themselves be archived off the machine.

Done when: The record accounts for its own closure; the archive contains the account.

4. What this placeholder honestly does not know yet

This draft has never been executed, and four questions await experience or research before v1.0: the exact mechanism for revoking trust grants in the current tool version (Step 6 — part of the Session 6 frozen-subject investigation's territory); whether the tool's stores contain locations this census has not yet met (the anatomy moved once during a two-day pilot; over a months-long matter it will move again); what a proportionate closure looks like for SMALL matters — this procedure is sized for the pilot's custody standard, and the break-even analysis warns that fixed governance floors price out the middle market, so a light-touch variant may be needed; and how privilege interacts with the transcripts — whether the AI's working record is disclosable, privileged work product, or something the profession has not yet categorised, which is a question for the SCL paper rather than a procedure.

5. The closing thought the pilot earned

Setup creates the known state; Wake-Up verifies it; Warm Close records it; Closure proves what remains of it. The four together are the custody envelope around a tool that — capably, honestly, and without ever being asked — spreads evidence into places nobody designated and remembers things nobody told it to keep. The envelope is not distrust of the tool; it is the professional answer to the question every procedure in this set exists to let the operator answer under oath: where is the evidence, and how do you know?