

Dossier Space Pilot — Anticipating Dissonance

Status: Design finding and working principle

Date: 4 August 2026

Authors: Ken Tombs with Leonard / ChatGPT (independent reviewer)

Related artefacts: Defensibility Scorecard Rubric v1.0; Meta-Primer v0.9; ADR-004 Custody Server; *Perfect Until It Is Not*; Session Planner v1.1

Purpose: Record an architectural tension identified before implementation proceeds, so the pilot anticipates dissonance rather than designing for permanent perfection.

1. Finding

The emerging Dossier Space architecture is intentionally conjunctive.

A defensible run must demonstrate, at the same time:

- **Procedural integrity (X)** — scope, provenance, custody, constraints and refusals;
- **ROMER discipline (Y)** — a visible and reconstructible path from evidence to result;
- **Agentic interrogability (Z)** — an independent record against which the AI's account can be challenged.

The Defensibility Scorecard therefore treats weakness on any one axis as material. Strength on one axis cannot simply compensate for failure on another.

However, the pilot has also established the governing principle:

All is perfect until it is not — then we want to know why, obviously and transparently.

Taken literally, these two ideas can appear to conflict.

The scorecard can be read as saying that a run falling below threshold is disqualified. The transparent-exception principle says that imperfect runs should remain visible, usable and open to human judgement rather than being automatically discarded.

This is not a drafting defect to be removed. It is a genuine architectural dissonance that the system must anticipate and manage.

2. The dissonance

Evidence work is not a sequence of perfectly resolved states.

In practice:

- stronger procedural controls may constrain analytical completion;
- a complete reasoning path may reveal that required evidence lies outside scope;
- an unresolved provenance issue may make a result less defensible while making it more evidentially important;
- a custody log may expose differences between what the AI reports and what it actually did;
- a hanging thread may be valuable for investigation while not yet suitable for release as a conclusion.

The three axes therefore do not merely measure independent qualities. They manage tensions between:

- capability and constraint;
- analytical completion and evidential uncertainty;
- imperfect work and the need for accountability.

The architecture must not assume that all three will always trade cleanly.
Its responsibility is to make unclean trades visible and attributable.

3. Resolution: separate defensibility from usability

The pilot should distinguish two decisions.

3.1 Defensibility gate

The Defensibility Scorecard answers:

Did this run demonstrate the predeclared evidential standard for this matter?

This remains a conjunctive assessment.

If one axis falls below the matter threshold, the run has **not demonstrated defensibility at that tier**.

A zero on an axis prevents the run's conclusions from being released as defensible output.

3.2 Human disposition

A named professional answers:

What may still properly be done with this run?

Possible dispositions include:

- accept;
- accept with qualification;
- return for further work;
- retain as an investigative lead;
- retain as a documented exception;

- release after authorised bypass;
- reject for the intended purpose.

A run may therefore fail the defensibility gate and still retain analytical, investigative or learning value. The framework does not make the final usability decision. It ensures that the imperfection cannot be overlooked when the human makes it.

4. The gate is an alarm, not a guillotine

The conjunctive gate should be interpreted as an **alarm**.

While all three axes meet the declared standard, the gate remains silent.

When an axis falls below threshold, the gate announces:

- which axis is affected;
- the criterion that was not met;
- the evidence for that assessment;
- any coupling with another axis;
- the effect on the result;
- whether human authority can release or qualify the result;
- what further action would be needed.

The gate does not delete the run, erase its observations or pretend the work never occurred. The filed record remains frozen and visible.

A failed or imperfect run may be especially important because it exposes the precise boundary, conflict or missing evidence that future work must address.

5. Practical rule for the pilot

The following working rule is adopted for future design and implementation:

A below-threshold run is not released as defensible at the declared matter tier, but it remains filed, visible, interrogable and available for an explicitly recorded human disposition.

This preserves both principles:

- the scorecard does not falsely certify imperfect work;
 - the architecture does not discard real evidential value merely because tension surfaced.
-

6. Required output when dissonance occurs

When the axes do not trade cleanly, the run record should contain an explicit exception block.

Suggested minimum form:

```
EXCEPTION / DISSONANCE RECORD
```

```
Run ID:
```

```
Matter ID:
```

```
Affected axis or axes:
```

```
Threshold not met:
```

```
Observed condition:
```

```
Why it matters:
```

```
Coupling effect, if any:
```

```
Effect on conclusion:
```

```
Defensibility status:
```

```
Permitted current use:
```

```
Human decision required:
```

```
Named decision-maker:
```

```
Further action:
```

```
Linked custody-log entries:
```

```
Linked ROMER elements:
```

Example:

```
Affected axes: X and Y
```

```
Observed condition: Relevant messages appear to exist outside the admitted custodian scope.
```

```
Coupling: Procedural compliance prevented analytical completion.
```

```
Effect on conclusion: Chronology remains incomplete.
```

```
Defensibility status: Not demonstrated for release at litigation tier.
```

```
Permitted current use: Investigative lead only.
```

```
Human decision required: Whether to expand matter scope.
```

7. Consequences for Sessions 8–10

Session 8 — custody-server implementation

The server should preserve and expose boundary events rather than merely block them silently.

A refusal or provenance divergence must create a citable log event.

Session 9 — measured mediated runs

The experimental record should capture not only performance and overhead, but also:

- how often the axes conflict;
- which axis combinations generate conflict;
- whether infrastructure reduces or merely relocates the conflict;
- whether the conflict is visible in the filed record;
- whether the run remains reconstructible after the conflict.

Session 10 — primer interaction

The primer-layering experiment should deliberately include cases designed to produce dissonance, including:

- out-of-scope but analytically relevant evidence;
- provenance divergence;
- incomplete evidence chains;
- a trace-versus-log discrepancy;
- a material analytical choice that cannot be fully justified;
- a bypass request requiring named human release.

Success is not that the primers make these tensions disappear.

Success is that the architecture identifies them, records them, prevents silent continuation, and presents them clearly for professional judgement.

8. Design warning

The pilot must resist two opposite errors.

Error 1 — perfection theatre

The system produces extensive green attestations and implies that compliance with the primers guarantees correctness.

This would recreate the exact governance problem Dossier Space is intended to solve.

Error 2 — automatic invalidation

The system treats every unresolved tension as fatal and discards the run.

This would make the framework unusable for real evidence work, where uncertainty, incomplete scope and conflicting duties are normal.

The correct position lies between them:

Do not claim perfection. Do not hide imperfection. Do not automatically destroy the value that remains.

9. Architectural interpretation

Dossier Space should be understood as an environment for **managed evidential dissonance**.

The axes are not expected to remain perfectly aligned at all times.

Instead:

- X exposes where procedure constrains capability;
- Y exposes where evidence does not support complete reasoning;
- Z ensures that the unresolved residue can be challenged;
- the Meta-Primer governs which boundaries cannot be crossed;
- the human decides what use, if any, may properly be made of the imperfect result.

This is not a compromise in the architecture.

It is the point at which the architecture becomes realistic.

10. Working conclusion

The purpose of Dossier Space is not to create a permanently perfect analytical environment.

Its purpose is to create an environment in which:

- clean work proceeds without unnecessary ceremony;
- tension is expected;
- exceptions become unmissable;
- the reason for the exception is recorded;
- no axis silently overrides another;
- the AI cannot release its own boundary crossings;
- imperfect work remains available for informed human judgement;
- the record shows both what held and what did not.

The architecture should therefore be designed around the expectation that dissonance will occur.

Perfect until it is not is not a promise of perfection. It is a promise that imperfection will announce itself.