Session 7 — Opening Note & Mid-Pilot Recap

Restart after a 17-day hiatus (15 July – 3 August 2026) · Ken Tombs with Claude (design lead, Opus 4.8) · This note opens the Session 7 log (Entries 082–089) and recaps the pilot's aims through to the introduction of dimensions and the scorecard. Prepared as the reconnection artefact after a genuine gap — itself an experimental condition, below.

Opening log entries

Entry 082 — Session 7 opened after a 17-day hiatus — recorded as an experimental condition

The pilot resumes after a genuine gap of seventeen days (last session 15 July; restart 3 August 2026) — far longer than any prior session boundary (previous boundaries were overnight or a lunch break). The hiatus is logged not as a scheduling fact but as a condition: it simulates the natural interruptions of real litigation — evidence preparation halted mid-matter, for example by a judge calling the parties to chambers, or a matter parked pending a ruling, or counsel pulled onto something urgent. Seventeen days is longer than many real adjournments, and it spans a period in which the tool would ordinarily have attempted several updates.

What it tests, that no prior session did: whether the custody architecture reconstructs a verified working state from RECORDS ALONE after a real gap, when neither operator nor design lead retains the details in working memory. This is the direct, live test of the pilot's founding doctrine — "the chat is a convenience; the document is custody." Until now that doctrine was asserted; the hiatus made the wake-up ritual prove it — and it did: every detail needed came from the disk and the banner, not from memory. The freeze verification gained a second dimension too: not merely "did the configuration hold overnight" but "did a frozen subject stay frozen across seventeen unattended days" — answered, with a precise and bounded result (Entries 084–089).

Entry 083 — Design-lead model confirmed: Opus 4.8 through Session 9, Fable 5 from Session 10

The design lead (the Claude instance directing method, drafting, forensics, and logging — not the experimental subject) runs on Opus 4.8 for Sessions 7–9, on cost-capability grounds: procedure design, forensic reasoning over pasted outputs, and document generation do not require frontier capability, and Opus economises the resource pool that heavy subject sessions also draw on — the pilot's own break-even logic applied to its own staffing.

Escalation to Fable 5 is reserved for Session 10 (the multi-dimensional primer / axiom / scorecard analysis — the densest reasoning of the arc, layering primers and scoring against independent logs) and for any implementation stretch (e.g. Session 8's server build) that demonstrably needs it, as a deliberate, logged act.

Doctrine preserved without exception: the design-lead model choice NEVER touches the experimental subject. Claude Code remains on Sonnet default throughout, per Entries 009/040 and the frozen-subject configuration deployed in Session 6. "Out of the box" means the unconfigured default; the subject is now literally pinned by settings.json. Whatever model the design lead runs, the subject is Sonnet.

Entry 084 — Store census across the hiatus: the tool's footprint frozen in time

The seventeen-day store diff: the tool's anatomy under ~\.claude unchanged (ten folders, no eleventh); every file and folder timestamped 14 July or earlier; the three session transcripts (102901e2, 7b34c04e, f50a4f7e) and the tool-results capture all present and frozen. Zero unattended activity across seventeen days. The machine that had been auto-updating every few days went completely silent the moment DISABLE_UPDATES was deployed and the channel dropped to stable. Operator's framing, adopted: stability is not a property the tool provides by default — it is an outcome the operator achieves by deliberately prioritising accuracy and consistency over currency. Sameness, like forgetting, must be engineered.

Entry 085 — A phantom anomaly, tested and dissolved: the "missing" outputs folder

pilot_outputs appeared absent from the initial -Force workspace listing; the operator provisionally attributed the loss to a prior operator error. Direct check confirmed the folder present and empty, exactly as left at Session 6 close — the initial reading was the error, not the folder; the provisional attribution retracted. The navigable-chain discipline catching itself: a graded "possibly" tested and dissolved, in the direction that matters — it stopped us manufacturing a human fault as readily as it would have stopped us missing one. Lesson: verify an anomaly by direct check before recording its cause.

Entry 086 — Locked-down structure as the fourth leg of the custody envelope

The restart surfaced three phantom anomalies in succession — a suspected duplicate transcript, a bench parked inside the corpus (PowerShell launched from an Explorer address bar rather than the Start menu), an apparently-missing outputs folder — every one resolved to operator misreading or launch-context accident, none to actual fault. The pattern indicates a fixed, locked-down file and folder structure would prevent this entire class of error: immutable geometry, permissions set read-stable, launch contexts constrained. Same philosophy as the mediation layer, one level down — enforce policy through structure, not vigilance. Recommendation: a dedicated workstation configured and locked to the task BEFORE work commences; configuration hardening as the fourth leg of the custody envelope, alongside primer, procedures, and mediation. We froze the subject; the natural next move is to freeze the workstation.

Entry 087 — The model roster shifted under the freeze — the freeze's exact boundary found

A plain launch brought the subject up on Opus 4.8, not Sonnet 5. Cause: over the hiatus the vendor's model roster changed on the server side (Opus 4.8 and Fable 5 arrived; Opus 4.7 gone) and the DEFAULT selection silently moved to Opus 4.8 — beyond the reach of any local configuration. Not operator error this time: the tool re-defaulted itself, and every launch since would silently have been Opus unless caught. The subject was manually re-pinned to Sonnet 5 (still offered, promo through 31 August — continuity preserved). Critical finding: DISABLE_UPDATES froze the client binary and the filesystem, but not the model roster or default — you cannot fully freeze a cloud-backed subject, because the models live on the vendor's servers, not in the pinned client. A frozen client is a stable interface to a shifting roster. The sovereign-inference argument gains its sharpest supporting fact.

Entry 088 — The freeze has limits; what is frozen must be verified every wake-up

Operator's governing principle, adopted: freezing or locking down the configuration and architecture has limits, and what is frozen must be VERIFIED at every wake-up, never assumed. The hiatus proved it both ways — the freeze worked where it could reach (client, filesystem) and failed where it could not (server-side roster). A lock-down is risk-reduction, not guarantee: it narrows what can drift, but the boundary of its control must be re-established every session by direct check. "We froze it, so it is still as we left it" is precisely the assumption the wake-up ritual exists to forbid. Standing rule: the model in use is now a mandatory banner-read at every launch — the one thing the freeze cannot hold. The banner-read is three checks, each mandatory and recorded: client version (freeze-held?), model in use (roster-drifted?), workspace path (launched-correctly?). This morning all three mattered: version held, model drifted, path was wrong twice.

Entry 089 — Freeze verdict — bounded stability, precisely mapped

Banner confirmed after re-pinning: v2.1.209, Sonnet 5. The complete, bounded verdict of the seventeen-day hiatus: client binary held EXACTLY at v2.1.209 (the previously-volatile tool did not update once); filesystem stores entirely untouched; model roster and default DRIFTED server-side and were corrected by hand. The freeze is real but bounded — it holds everything local and nothing server-side. Verification, not the freeze, is what guarantees the subject's identity at any given session. Stronger for containing its own limit: not "we achieved perfect stability" (which invites the first counterexample) but "we achieved bounded stability, mapped its exact boundary, and built verification to cover the remainder." The wake-up ritual reconstructed a fully verified state from records alone across a two-week gap, caught the one thing that drifted, corrected it, and recorded the boundary — the founding doctrine demonstrated, no longer merely asserted.

Recap — the pilot's aims through to dimensions and the scorecard

The pilot asks a single question in stages: can agentic AI be trusted with legal evidence, and what makes it trustworthy? Sessions 1–9 answer the infrastructure half — capability, governance, custody, and mediation. Session 10 opens the method half — the primer apparatus itself. The table records each stage's aim and where it stands at this restart.

Session	Aim	Status / outcome
1–5 (complete)	Establish and close the comparison design: can out-of-the-box AI do evidence work, and what does a primer add?	DONE. Retrieval verified 9/9 precision and recall against independent ground truth under three conditions (primed/4.6, baseline/5, primed/5). Capability constant; the primer's contribution isolated as exhibit discipline — accurate work becomes citable work. Governance overhead measured (≈ 3× approvals under the primer). 70 entries.
6 (complete)	Harden the boundaries: wake-up and sleep as documented procedures; recover what was lost; freeze the subject.	DONE. Four procedures drafted and two executed; Session 5 fully harvested; the hidden .git solved (custody instrument broken by the rename it might have recorded); the destroyed report recovered from the tool's checkpoint store; the frozen-subject configuration deployed — and the volatility's root cause found in a settings file ("autoUpdatesChannel": "latest"). 11 entries (071–081).
7 (opening)	Design the laboratory bench on paper: ADR-004 — the MCP custody server, frozen before any code.	Tool set (read_file, list_folder, search_messages, get_headers, hash_document, verify_hash); self-describing definitions; scoping and privilege-gate rules; chained custody-log schema. Side-deliverable: the Session 10 scorecard rubric, fixed in advance.
8	Build and smoke-test the server; connect Claude Code in passthrough mode (Phase 2a).	Working server; first mediated run; plumbing proven. Design-lead escalation to Fable 5 available if the code strains.
9	Measured Phase 2 runs: 2a passthrough vs 2b custody-on, against the existing grid.	ADR-003's central number — the cost of governance, isolated. The trace-versus-log gap measured for the first time (an infrastructure log finally exists).

Session	Aim	Status / outcome
10 (the turn)	The multi-dimensional analysis: multiple primers in concert, axioms as a distinct stratum, and the scorecard applied.	Where governance-as-method meets governance-as-infrastructure. Design lead escalates to Fable 5. Beyond this point the pilot measures the primer apparatus itself, scored against Session 9’s independent logs.

The through-line, in one paragraph

The pilot has established that a stock AI retrieves evidence with verified accuracy, that a primer converts accurate work into citable work (exhibit discipline), and that the real risks live not in capability but in accountability — memory that persists invisibly, versions that drift silently, authority that blurs, and records that must be built rather than assumed. Sessions 7–9 answer these with infrastructure: a mediation layer (the MCP custody server) that enforces policy structurally and generates a custody record independent of anything the AI reports about itself, measured against the Phase 1 baseline. Session 10 then turns from infrastructure to method — introducing multiple primers in concert, axioms as a distinct governing stratum, and a scorecard that makes "exhibit discipline" a measured property scored against Session 9’s independent logs rather than a characterisation. Everything up to Session 10 builds the bench; Session 10 begins to measure what the analyst does at it.

Immediate agenda for Session 7

The wake-up ritual is COMPLETE (Entries 084–089): it reconstructed a verified state from records alone across a seventeen-day hiatus, dissolved four phantom anomalies, caught one real drift (the model default), corrected it, and mapped the freeze’s exact boundary. The founding doctrine — document is custody, not memory — is demonstrated. The subject is re-pinned to Sonnet 5 and, this being a bench-and-drafting session, exited; no measured subject runs today. The session now turns to its substance: ADR-004, the MCP custody server designed and frozen on paper (six tools; self-describing definitions; scoping and privilege-gate rules; the chained custody-log schema), with the Session 10 scorecard rubric drafted alongside it. Brief housekeeping also outstanding from Session 6: deduplicate GIT Deliverables; locate the mislaid break-even note; reissue the procedure v1.0s with Session 6’s and today’s earned amendments (notably: Start-menu launch for the bench; the three-part mandatory banner-read; workstation lock-down at setup).