

Glossary — the Pilot's Working Vocabulary

v1.0 · 15 July 2026 · Terms coined or specialised during Sessions 1–6, defined for collaborators, assurers, and readers of the published set. Where a term encodes a lesson, the lesson is stated — the vocabulary is itself part of the record.

The subject	The AI under test: Claude Code running on the pilot workstation, out of the box, whose behaviour is the research data. Per ADR-003 Layer 1: never modified, extended, or compensated for. Redefined after the July auto-updates as "Claude Code current", versions recorded at every launch.
The bench / audit bench	A plain PowerShell window with no AI in it, used for all verification and custody operations. Retitled "AUDIT BENCH", self-transcribing via Start-Transcript, never parked inside the workspace. The bench exists so that no check is ever executed by the thing being checked.
The design lead	The Claude instance in the chat interface (this collaboration) directing method, drafting documents, and keeping the log. Has no access to the workstation; everything it knows arrives through the operator. Same model family as the subject — independence lives in the architecture, not the parties.
The operator	The human (Ken) executing every command, holding authority at every approval gate, and owning the record. The only channel between subject, bench, and design lead.
Workspace	The single folder the subject is trusted on and launched from (enron_mail_20150507). Everything the subject may see or touch lives here: the deployed primer, the evidence (read-only), and pilot_outputs.
Exhibit level	The folder one level above the workspace (ENRON Data Set Test 1) — outside the subject's trust boundary, therefore invisible to it. Home of run exhibits: archived memories, transcripts, per-run outputs, banked primers. Safe by geometry so long as the subject is only ever launched from the workspace.
GIT Deliverables	The folder at the pilot top level holding the document set (logs, procedures, synthesis, planner) destined for public GitHub publication. Documents ABOUT the pilot live here; evidence FROM the pilot stays at the exhibit level.
Primer	A standing instruction document governing the subject's conduct. Two in service: the six-line safety primer (constraints only — the baseline condition) and CLAUDE.md v0.1 (the full standing primer: CIPDA stages, ROMER outputs, citation discipline — the primed condition / Phase 3a treatment). Deployed by copying into the workspace as CLAUDE.md; identity verified by byte size, never impression.
CIPDA	Context, Intent, Plan, Deliver, Assure — the principal investigator's structural extension of PDCA for hybrid human-AI reasoning. Stage primers are CIPDA-structured.
ROMER	Reasoning, Observations, Methods, Evidence, Results — the output structure for evidential reporting. In Run 1.2-P5 the subject assembled it by name, unprompted, from the primer's prose description alone.
Warm / cold	Session states. Cold: the subject demonstrably knows nothing beyond what the deployed primer discloses. Warm: prior context present (same session, resume, or memory). Doctrine: cold must be verified, never assumed; warm is legitimate only as a recorded condition, never as an accident.
Wake-up ritual	The start-of-session procedure verifying machine state against the previous close before the subject is touched: bench first, store census, workspace check, primer identity, versions at launch, primer-loaded question, coldness check.

Warm close	The end-of-day procedure writing the record tomorrow verifies against: work finished or formally abandoned, outputs harvested, shells killed, resume ID captured at exit, transcript archived same-night, closing state written, bench closed last. Distinct from Final Closure.
Filed means frozen	Doctrine: an output saved to disk is a fixed artefact. Corrections go in subsequent outputs or conversation, never edits to filed documents. Adopted after the tool twice treated filed evidence as drafts to improve.
Forward-only corrections	Wrong inferences stay in the record beside what corrected them; no entry is rewritten. An audit trail that only records successes is not an audit trail.
Session / run / entry	Three clocks, never conflated: sessions are bounded by operator absence; runs by subject launches; entries by findings.
Session Zero	The first-time initialization: the recorded birth state (folder geometry, trust grants, store anatomy, primer deployment) that every future wake-up verifies against.
Ground truth	An answer set constructed entirely outside the subject (bench commands against the corpus), against which the subject's retrieval results are scored. The pilot's: the nine lay-k December 2001 "bankruptcy" messages.
The trust boundary	The folder scope Claude Code has been granted read/edit/execute over via its trust prompt. Grants persist indefinitely and invisibly; the pilot's geometry rule exists because of this.
The stores / footprint	The tool's own record locations under the user profile (~\.claude): projects (transcripts, tool-results), sessions, file-history (per-file checkpoints), backups, paste-cache, session-env, shell-snapshots — an anatomy that moves between versions. Collectively, where evidence-derived fragments accumulate without being asked to; threat when unmapped, safety net when mapped (the destroyed report was recovered from file-history).
Frozen subject	The configuration deployed in Session 6 pinning the subject's behaviour: updates blocked (DISABLE_UPDATES), auto memory disabled, prompt suggestions disabled, update channel dropped to stable. Set-then-verify: each switch is confirmed behaviourally before procedures rely on it.
Phase 1 / 2 / 3	ADR-003's phases. 1: baseline, direct filesystem access. 2a/2b: MCP-mediated — passthrough, then custody-on (hashing, chained logging, scoping). 3a/3b: methodological governance — primer, then primer plus ROMER templates.
MCP	Model Context Protocol — the open standard by which an AI client uses external tools. Phase 2's custody server is an MCP server: Leonard's "laboratory bench" — policy, not intelligence.
The comparison grid	The pilot's central result table: the same retrieval task at 9/9 under primed/4.6, baseline/5, and primed/5 — capability constant, exhibit discipline tracking the primer, governance overhead tracked as approval counts.
Exhibit discipline	The qualities that make a report citable rather than merely accurate: unique identifiers, calibrated confidence, scope fencing, absence-flagging, attribution. What the primer reliably restores.
Verification theatre	A check performed in form but not substance — the Session 1 read-back accepted on impression while the wrong primer sat on disk. The named failure mode the byte-level rule exists to prevent.
Authority channel	The subject's input line — where operator intent enters the system. Named after the discovery that the tool drafts prompts into it (grey ghost text), making "who

	instructed this?" a three-answer question: operator, nobody, or the tool ratified by a keystroke.
--	---