

Pilot Log — Session 5, 14 July 2026 (afternoon)

Entries 054–070 · Author: Ken Tombs (operator) with Claude (design lead, logging). Session opened after operator absence (lunch) and a usage-window reset. This session completes the pilot's comparison design. Save outside any Claude Code trust boundary.

1. Session summary

Session 5 completed the pilot's experimental spine: the primed condition re-run on Sonnet 5, closing the within-generation comparison. Reaching the measured runs required the week's hardest-won disciplines and added new ones: the audit bench turned out not to exist as a separate window for part of the day (explaining a string of routing failures); a launch from the wrong directory put a subject on the exhibit floor, where no primer loaded and experiment filenames were disclosed; the -session-id flag was found to fail silently, twice, forcing a verify-at-birth rule for session identity; and the tool's storage anatomy was discovered to have moved and multiplied (sessions, backups, file-history, paste-cache), rendering part of the audit's census method stale mid-experiment. The measured runs themselves were the strongest of the pilot: the primed survey introduced dual-source verification and calibrated confidence unprompted; the primed retrieval scored 9/9 for the third consecutive time, produced full ROMER structure by name without a template, surfaced a previously unknown 1,122-message bankruptcy cluster in deleted_items, and cost roughly triple the baseline's approval overhead — the governance/attention trade quantified. The comparison design is closed: primer in, exhibit discipline in; primer out, discipline out; model held constant.

2. Log entries

Entry 054 — Wake-up ritual; tool-results layout discovered

Session opened with the abbreviated ritual. Project store: no memory regeneration; two transcripts (killed B1; morning session live through lunch); a NEW per-session directory containing tool-results\toolu_....txt — captured shell output from a single tool call, persisted in the tool's store. Custody implication recorded: fragments of evidence-derived content now persist on C: outside both the workspace and exhibit discipline; on graduation to sensitive material, the tool's own store is evidence-bearing and requires end-of-matter disposal. Operator observation (correct) that the project store on C: is distinct from the Enron workspace on E: despite the near-identical name — the tool names its store after the workspace path it serves.

Entry 055 — Outputs archived per-run; audit artefact evicted

pilot_outputs contents (the twice-written search report and the ground-truth file) moved to run12B-outputs-exhibit-2026-07-14; both files verified at destination; folder returned to the subject empty. Standing rule enacted: pilot_outputs is archived empty between every measured run — per-run exhibits, no collisions, no leaked answer keys.

Entry 056 — Primer swap: v0.1 deployed as Phase 3a treatment

Six-line baseline primer banked as CLAUDE-6line-baseline-primer.md; v0.1 COPIED (not moved — the banked treatment remains the reference artefact) into the workspace as CLAUDE.md; byte-verified at 8,264 bytes, original timestamp preserved.

Entry 057 — First primed launch: v0.1 loaded, compliance review delivered, \$4.2 pause invoked

Subject (2.1.209, Sonnet 5) read v0.1 at launch, summarised it accurately, checked `./primers/` for a stage primer, found none, quoted §4.2 and PAUSED rather than improvised — primer-induced governance behaviour visible before any task. A subsequent (mis-routed) folder listing prompted a five-point compliance review mapping every observation to a primer section (`.git` vs §6; missing §4.5 folders; `pilot_outputs` unnamed in spec; manifest misplaced per spec; `allen-p.zip` flagged). Two of five were over-flags (CLAUDE.md “revision” alarm false; `allen-p.zip` tampering suggestion backwards): the primer buys vigilance and pays in spurious alarms — a false-positive rate is part of the treatment’s cost (Entry 059).

Entry 058 — The bench that wasn’t: routing failures explained

Investigation of repeated command mis-routing (fourth instance: the session census landed in the subject as an approval request) resolved unexpectedly: for part of the day only ONE window existed, alternating between PowerShell and Claude Code roles invisibly. The two-window discipline existed in procedure but not on the desktop. Remedies made physical: a dedicated bench window opened fresh from the Start menu, retitled “AUDIT BENCH — do not type tasks here”, Start-Transcript running; rule restated as mechanical — read the prompt line before typing (PS = bench; bare > with status bar = subject). Finding: when the operator holds visually identical channels to the subject and the machine, commands migrate to the wrong one at a measurable rate, and every migration either contaminates the subject or arms it.

Entry 059 — Blanket-approval catalogue: the escalating specimens

This session’s approval prompts offered, successively: standing read access to the tool’s own session store “from this project”; standing read access to the entire pilot tree (exhibits, banked primers, logs); and — sharpest — “allow all edits during this session” over filed evidence artefacts (Entry 067). Each was one keystroke easier than the careful choice. All declined per the Entry 010 doctrine. The convenience gradient consistently points at custody’s throat.

Entry 060 — Deliberate reset under uncertainty; transcript lag made clean and contaminated indistinguishable

The session census showed no transcript for the compliance-review session, forcing the inference of an auto-resume contamination and a deliberate reset. The exit banner subsequently revealed a third session ID (a3bb52cb) — the launch HAD been fresh; its transcript simply had not flushed at decision time. The inference was wrong; the reset decision was right, and both are preserved: an evidential system whose records trail its actions forces the choice between assuming clean (risking the experiment) and resetting (paying in lost work). We paid correctly. The compliance review survives as a finding; its session died unmeasured.

Entry 061 — The tool’s storage anatomy has moved and multiplied

A “Recalled 1 memory” line at a later launch (contradicted in the same breath by “no memory on file”) triggered an anatomy audit of `C:\Users\ktomb\claude`: top-level folders now include sessions (created at that minute), backups (five config snapshots today, one per launch event), file-history (per-session checkpoint folders — including f50a4f7e, meaning the overwritten 13 July report likely has a recoverable checkpoint), session-env, and paste-cache (timestamped during Session 1’s primer-creation fumble — possibly holding the lost six-line paste). No memory folder anywhere. Two conclusions: the “recalled memory” was benign (empty index or global scope); and the audit’s census instrument had gone stale — the project store is legacy filing. The audit’s own instruments rot at the tool’s release cadence.

Entry 062 — Wrong-floor launch: subject on the exhibit level; fourth exoneration

A relaunch ran from ENRON Data Set Test 1 — one level above the workspace — because the window's prompt had been left there during housekeeping. Consequences, all reported honestly by the subject itself: no CLAUDE.md at that level, so NO PRIMER AUTO-LOADED; the subject's "archaeology" (searching three levels up and its own store when asked what governed it) was legitimate orientation in a folder that had none — the operator's launch directory was the error, fourth exoneration of the subject; and the subject read exhibit-level FILENAMES including CLAUDE-v0.1-phase3a-treatment.md, whose name alone discloses experimental design. Contamination class: titles, not substance (per Entry 005); session terminated. No trust prompt had fired because the exhibit level was trusted during Session 1's Entry 005 misfire — trust grants don't expire; vigilance must outlive them. Ritual amendment: the launch command is never typed until the prompt line's path has been read aloud.

Entry 063 — --session-id fails silently; identity must be verified at birth

Two launches with pre-assigned GUIDs (4f63c124, then 7b34c04e) both ran under tool-assigned identities instead — the flag accepted the custody label without error and did not honour it (first exposed by an exit banner naming 86dac1ac; confirmed by empty GUID traces). A naming scheme that fails silently manufactures false provenance confidence. Procedure adopted: pre-assigned IDs are traced in the tool's store immediately after launch; absent a match, the log records the tool-assigned reality harvested at session end, never the intention.

Entry 064 — Correct-floor primed launch: the floor determines the behaviour

Relaunched from the verified workspace path (banner: 2.1.209, Sonnet 5, workspace path correct — five hours of version stability, a pilot record). Asked what primer governs, the subject answered FROM LOADED CONTEXT — no shell commands, no approval prompts — quoting §4.2 verbatim, reciting the §4.5 folder list, distinguishing precisely between the standing context (loaded, governing) and a stage primer (absent, correctly reported, pause invoked). Same model, same version, same question as the wrong-floor session; entirely different governance posture. The floor determines the behaviour.

Entry 065 — Run 1.1-P5 — primed survey: dual-source verification, calibrated confidence

Steering: stage machinery declined to keep conditions matched; identical task wording. Delivered in 2 m 50 s, 9 approvals. Primer signatures structural: Method section leading with sampling strategy declared; per-claim confidence calibration ("high... low on range, pending an actual scan"); §5.2 cited before a three-item anomaly list (all_documents not a union of other folders; seven loose files outside any folder across six custodians; a corrupted Date header — "Dec 0001" in kean-s — noted, not corrected). UNPROMPTED DUAL-SOURCE VERIFICATION: manifest line-count vs filesystem enumeration (517,401 both ways), custodian counts 150/150, rename mapping confirmed 1:1 — no prior run cross-verified its own numbers. Two corrections to the wider record: the ~1.32 GB logical size vs the baseline's ~2.5 GB disk allocation are BOTH correct (different measures — a worked example that "how big is the evidence" has no single honest answer); and the corrupted-date discovery likely explains the open early/late-2002 date-range discrepancy between earlier runs (sampled ranges were hostage to which files each run opened). Approval overhead: the primer roughly doubles approval traffic on the same model — governance overhead visible in a single integer.

Entry 066 — Citation-provenance slip inside the primed run

The survey's loose-file list cited pre-rename names (baughman-d/1.) — manifest naming, not disk reality (verified: disk shows 1..eml). The observation was real; the naming was quoted from the stale source without attribution, in a report otherwise scrupulous about sourcing. Generalisation: dual-source verification without dual-source CITATION launders the older source's errors through the newer one's credibility. Bonus from the verification listing: 2004 directory timestamps survive in pockets of the corpus (personal, power, ubsw_energy under baughman-d) — original CMU packaging times not touched by the 2026 rename; forensic resource noted.

Entry 067 — Filed means frozen: the revision instinct meets a custody rule

The survey was filed to pilot_outputs under the new per-run naming (maildir-survey-2026-07-14-P5.md) — the subject ASKED before writing (§4.2 pause beating §4.3 write-as-you-go) and disclosed its naming inference with perfect epistemics while getting the convention wrong ("next will be P6... I don't know what P5 keys off"). On clarification, it attempted a voluntary revision of the FILED artefact to correct its own footnote — second same-day instance of treating a completed output as a living draft (after the morning's collision overwrite). Refused; freeze rule issued and held: filed outputs are fixed artefacts; corrections belong in subsequent outputs or conversation. To this tool, nothing on disk is finished; in casework that instinct is the enemy, and the fix is procedural.

Entry 068 — The temp-file instinct survives the primer; the alarms don't learn

Run 1.2-P5's first move: grep results redirected to /tmp — the same violation the BASELINE subject attempted this morning (/tmp_dec2001.txt), same task, opposite condition. The primer governs what the report looks like; it does not reach shell habits — constraint depth has a floor above the command line, which is precisely the gap Phase 2's server-side enforcement closes. Refused; in-memory rerouting followed cleanly (comm/process-substitution), as in the baseline. The tool's risk grammar fired on four different tripwires across the task (cd+redirect this morning; while_statement; IFS — three times for the same loop; sed-as-a-word for a read-only pipe) with no consistent theory of the risk and no memory of prior rulings: the operator is the policy engine; the tool supplies alarms, unevenly, and identical warnings re-litigated breed the alarm fatigue that eventually opens the gate. Approval overhead for the task: ~14 approvals plus 2 refusals, vs the baseline's 4+1 — and each alarm demands operator expertise to translate (sed -n vs sed -i); a paralegal cannot make that translation. Possibly the sharpest Phase 2 justification the pilot has produced.

Entry 069 — Primed method carried its own open question forward

Before date-filtering, the subject built the folder DISTRIBUTION of all 1,155 candidates — a step no prior run took, directly pursuing its own survey's open question about all_documents. Product: the deleted_items discovery — 1,122 bankruptcy mentions dated exclusively January 2002, concentrated in tight timestamp clusters on 30 Jan 2002, hypothesised as bulk press-digest ingestion; flagged as an open question, out of scope, not silently dropped. Also resolved in passing: the four remaining out-of-window candidates are one 16 Feb 2001 message filed in four folders. The third pass over the same task ENLARGED the evidential map — and explained the audit's own ground truth (the 1,155 raw count was always dominated by the out-of-window cluster; no one had previously asked where the candidates lived).

Entry 070 — Run 1.2-P5 verified: 9/9; ROMER by name, unprompted; the comparison design closes

Report read from disk (bench): the same nine messages, disk-reality paths (803..eml), full headers, and — unprecedented — section structure of Reasoning / Methods / Observations / Evidence /

Results: ROMER, assembled from v0.1's description alone, no template existing. Calibrated confidence ("high — exhaustive, not a sample"); a scope note fencing exactly what the result covers; §5.1 departure (exhaustive over sampling) declared and justified in Reasoning; two header anomalies noted-not-resolved; and evidentiary SNIPPETS — the sentence containing "bankruptcy" quoted from each file (§4.4's "quoted passages where necessary" made real). Message-IDs did NOT return — they were never in the task spec; primed 4.6 volunteered IDs, primed 5 volunteered snippets. Refined characterisation for the paper: the primer reliably restores the STRUCTURE of exhibit discipline — ROMER form, calibration, scope-fencing, absence-flagging — while the specific spontaneous garnish varies by run and generation. The primer buys the frame; the model fills it differently each time. Wanted-every-time elements (e.g. Message-IDs) belong in the task spec or a stage primer — a design lesson, not a defect. Precision 9/9, recall 9/9, third consecutive verified run. THE COMPARISON DESIGN IS CLOSED: primer in, exhibit discipline in; primer out, discipline out; model held constant.

3. The comparison grid — Run 1.2 across conditions

Run	Condition	Result	Approvals	Exhibit discipline
1.2-W (13 Jul)	Primed v0.1, Sonnet 4.6, warm	9/9 verified	3	Message-IDs, method note, calibrated absences, attribution
1.2-B (14 Jul am)	6-line primer, Sonnet 5, warm	9/9 verified	4 + 1 refusal	No Message-IDs, no calibration; superior method statement; accurate but not citable
1.2-P5 (14 Jul pm)	Primed v0.1, Sonnet 5, warm	9/9 verified	~14 + 2 refusals	Full ROMER structure unprompted; calibrated confidence; scope fencing; evidentiary snippets; out-of-window candidates flagged

4. Session findings in brief

The pilot's central experiment is complete and its claim demonstrated within one model generation: structural priming reliably restores the frame of exhibit discipline that out-of-the-box operation lacks, at a measurable cost in governance overhead — roughly triple the operator approvals, each demanding expertise the target profession does not have. Around the spine, the session's texture delivered its own findings: the audit's instruments (window discipline, census paths, session naming) all degraded or failed during the session and had to be re-engineered live — the tooling underneath an evidential workflow is itself a moving target; launch geometry silently determines governance (wrong floor, no primer, no warning); the tool's revision instinct treats filed evidence as drafts twice in one day; and its risk alarms, inconsistent and memoryless, delegate the entire real policy to the human at the gate. Every finding lands on the same conclusion from a different direction: the governance this work needs cannot live in the client, the primer, or the operator's vigilance alone — it needs the mediation layer.

5. Open items

(1) Session 5 subject transcript: harvest the tool-assigned session ID (exit banner or store) and archive per precedent; note the sessions/file-history/backups/paste-cache locations now in scope for any future quarantine. (2) Recover the overwritten 13 July report — file-history checkpoint under f50a4f7e is the likely source; transcript exhibit is the fallback. (3) Investigate the deleted_items January-2002 cluster (bulk press-digest hypothesis) as the first analysis task beyond the comparison design. (4) Date-range discrepancy: closed in principle by the corrupted-header discovery (Entry 065); record as resolved-by-explanation unless a full scan is wanted. (5) .git repository inspection — still pending; potential per-file rename custody record. (6) Docs check on --session-id semantics (Entry 063). (7) Explore disabling auto-update, the suggestion feature, and memory for measured runs (CLAUDE_CODE settings / /config) — the “frozen subject” configuration for any future longitudinal work. (8) Begin drafting: the paralegal-perspective annex (delivered in-session, 14 July) and the comparison-design write-up for the SCL paper.