

Pilot Log — Session 6, 15 July 2026

Entries 071–079 · Author: Ken Tombs (operator) with Claude (design lead, logging). Session scope per the Planner: boundaries — procedures executed, recoveries performed, the subject frozen. No subject launch occurred; a wholly bench-work session. Save to GIT Deliverables.

1. Session summary

Session 6 executed the Wake-Up Procedure v0.9 for the first time (passing, with amendments captured), reconstructed and completed Session 5's missing close (transcript located and archived, outputs exhibited, two standing mysteries resolved by a single timestamp), solved the hidden .git repository (finding the custody instrument broken by the very rename it might have recorded, plus a line-normalisation hazard that made the abandoned full load a lucky escape), recovered the pilot's one destroyed artefact from the tool's file-history checkpoints, and deployed the frozen-subject configuration — discovering en route that the pilot's notorious version volatility had been a settings choice ("autoUpdatesChannel": "latest") sitting in the tool's config all along. Earlier in the day, before the session opened, the four operating procedures and the Session Planner were drafted; morning decisions also fixed the deliverables' home (GIT Deliverables, destined for public GitHub with external assurance by Leonard and Mistral) and the division of labour (design lead runs the pilot; assurance and publication are external).

2. Log entries

Entry 071 — Session opened; Wake-Up v0.9 first execution begun

Session 6 opened by executing its own deliverable: the Wake-Up Procedure draft, end-to-end, timed. Steps 1–3 (log first, bench built and self-transcribing, parked at exhibit level) passed cleanly.

Entry 072 — Store anatomy grew again: shell-snapshots

Step 4's census found a TENTH top-level folder — shell-snapshots, created 14/07 15:35, mid-Session-5 — against the nine-folder baseline recorded a day earlier. Entry 061's warning (the audit's instruments rot at the tool's release cadence) validated by the procedure built on it. v1.0 baseline list amended.

Entry 073 — The 16:59 "oddity": resolved by the missing close

Operator noted store writes at 16:59 on 14/07, after work had (probably) finished; machine recollected as cold immediately after. Resolved when the Session 5 transcript surfaced with LastWriteTime 16:59:52: the primed session was never exited and died with the machine's shutdown — the writes were its termination flush. Unresolvable at first only because no warm close had run; the closing-state record that would have answered it did not exist. The Warm Close procedure justified by the gap it will close. Standing lesson: activity windows are defined by session lifetime, not operator attention.

Entry 074 — --session-id vindicated; Session 5 transcript harvested

The census found file-history\7b34c04e... — the operator-assigned GUID from Session 5, which the pilot had written off as a silent flag failure (Entry 063). The flag had worked; the TRACE had failed — searched too early and in the wrong places, with the wrong-floor session's exit banner reinforcing the misreading. Fifth instance of the pattern: when tool and auditor disagree, check the auditor first.

Session artefacts confirmed across three stores (file-history, projects, session-env). The transcript (294,945 B, terminated-by-shutdown) archived to session5-transcript-exhibit-2026-07-14, size-verified. Entry 063's verify-at-birth rule stands, with corrected search locations; the flag returns to service.

Entry 075 — Wake-Up execution completed; Session 5 books balanced

Steps 5–6 passed: workspace exact, CLAUDE.md at 8,264 B (v0.1 by identity). The two P5 outputs moved to session5-P5-outputs-exhibit-2026-07-14, verified both ends; pilot_outputs empty. Steps 7–10 executed-by-omission: no subject launch this session. Amendments captured for v1.0: merge Steps 5/6 (one -Force listing serves both); add shell-snapshots to the Step 4 baseline; add a branch point after Step 6 — no-launch sessions proceed directly to tasks, Steps 7–10 apply only before the first launch.

Entry 076 — The .git repository: broken instrument, modest treasure, dodged hazard

Forensic method: hooks inspected first (all .sample — inert); the repository copied whole to git-repo-exhibit-2026-07-15; the ORIGINAL untouched throughout. Diagnosis: git reported "not a repository" because the 3 July rename script had appended .eml to git's extension-less vital files (HEAD.eml, config.eml, index.eml) — the custody instrument damaged by the very event it might have recorded, and sitting silently broken for twelve days until an audit looked. Structure dates revealed the repo predates the rename by seven weeks (born 12 May, at corpus assembly).

The copy was repaired (suffixes stripped) and interrogated: ONE commit — 9d6297d, "Initial load / From CMU May 7, 2015 tarball", authored XPlain-R portal-admin, 12 May 2026 16:56 — tracking a single file, .gitattributes. The hoped-for per-file custody record does not exist; the full load was never committed. Recovered instead: a contemporaneous, authored, dated declaration of corpus source at assembly — provenance the record lacked, now cryptographically anchored.

The sting: .gitattributes contains "* text=auto" — automatic line-ending normalisation. Had the 517,401-file load ever completed, git would have stored LF-normalised versions differing byte-for-byte from the evidence. The abandoned load was a lucky escape: version control is not evidence-neutral by default. SCL keeper.

Entry 077 — The destroyed 13 July report: recovered

file-history\f50a4f7e... held one checkpoint (7d18466d3222feb6@v1, 6,542 B, stamped 13/07 17:05:52 — the original save moment). Contents verified by fingerprint (Generated: 2026-07-13, the PowerShell method line). Copied to recovered-artefacts-exhibit-2026-07-15 as RECOVERED-search-lay-k-bankruptcy-dec2001-2026-07-13.md; 6,542 B verified, original timestamp carried. Provenance honestly stated: recovered from the tool's checkpoint store, not the destroyed path. The finding beneath: the same invisible persistence that made forgetting hard made recovery possible — the stores are threat when unmapped, safety net when mapped.

Entry 078 — Frozen-subject investigation: all three switches exist

Documentation research (current, not recalled): auto-update — DISABLE_AUTOUPDATER stops the background check; DISABLE_UPDATES blocks all paths including manual (the true freeze); autoUpdatesChannel stable/latest plus minimumVersion for channel pinning; documented inconsistencies between docs and behaviour noted. Auto memory — the Session 2 episode is a documented feature; CLAUDE_CODE_DISABLE_AUTO_MEMORY=1 disables; field reports of the opt-out failing in some versions noted. Prompt suggestions — /config toggle or

CLAUDE_CODE_ENABLE_PROMPT_SUGGESTION=false (per the tool's own creator). Bonus for Setup v1.0: CLAUDE_CODE_DISABLE_NONESSENTIAL_TRAFFIC bundles updater/bug/error/telemetry switches — narrows outbound surface to inference alone; relevant to the sovereignty chapter. Also noted: a transcript-retention setting (default 30 days; 0 disables session persistence entirely) — a two-faced lever: kills auto-resume at the root but also the transcripts the pilot harvests as evidence; kept ON here, option recorded for privacy-constrained deployments. Doctrine applied to all: set, then verify behaviourally.

Entry 079 — Frozen-subject configuration deployed; root cause of volatility found

Rationale on the record: three findings trace to uncontrolled tool behaviours — version drift (039/046), unrequested memory (027), tool-drafted prompts (053); switches exist for all three; deployed at user level; each to be verified before procedures rely on it.

The existing settings.json read before editing — and contained "autoUpdatesChannel": "latest": the subject had been configured to take every release the moment it shipped. The pilot's notorious volatility was partly a settings choice, not merely vendor behaviour; "stable" was available all along. Also present: "theme": "dark-daltonized" — which sharpens Entry 053: the suggestion feature's only authority cue is a font colour, on a terminal explicitly configured for a colour-impaired operator. SCL keeper.

Merged file deployed and read back word for word: channel dropped to stable (fallback if the block ever fails), theme preserved, env block added — DISABLE_UPDATES=1 (total block chosen over background-only: unfreezing must be a deliberate, logged edit), CLAUDE_CODE_DISABLE_AUTO_MEMORY=1, CLAUDE_CODE_ENABLE_PROMPT_SUGGESTION=false. Empirical verification scheduled for the next subject launch: banner still 2.1.209, no memory writes, no ghost text.

3. Documents produced this day

Before the session: Session Planner v1.0; Wake-Up Procedure v0.9 DRAFT; First-Time Setup Procedure v0.9 DRAFT; End-of-Day Warm Close Procedure v0.9 DRAFT; Final Matter Closure Procedure v0.5 PLACEHOLDER. During: Glossary v1.0; this log. Standing decisions recorded: deliverables home is GIT Deliverables (public GitHub; external assurance by Leonard and Mistral); design lead's single focus is the pilot — publication legwork is external.

4. Closing state (Warm Close Step 6 record)

Subject: never launched this session; last known 2.1.209 / Sonnet 5; frozen-subject configuration now deployed, unverified. Workspace: v0.1 deployed (8,264 B); pilot_outputs EMPTY; corpus untouched (zero subject writes across six sessions). Exhibits created today: session5-transcript-exhibit-2026-07-14; session5-P5-outputs-exhibit-2026-07-14; git-repo-exhibit-2026-07-15 (repaired copy; original .git untouched in workspace); recovered-artefacts-exhibit-2026-07-15. Tool stores: ten top-level folders (shell-snapshots new); no memory folder. Open items carried: verify the three freeze switches at next launch (Session 7 opening act); issue procedure v1.0s incorporating today's amendments; deleted_items lead, date-range closure note, and all Planner items stand. Expected overnight: nothing — for the first time, the subject is configured not to change; tomorrow tests that claim.

5. The session's shape, for the record

A session in which the subject never woke, and which may still be the pilot's most productive: every mystery on the books was closed by bench forensics (the 16:59 writes, the --session-id "failure", the unreadable repository, the lost report), the procedures survived first contact and improved, and the volatility that shaped half the experimental design was traced to two lines of configuration. The boundaries session proved the boundary thesis: the machine's edges — start-up, shut-down, the stores, the settings — are where the evidence lives.