

Session 7 — Opening Note & Mid-Pilot Recap

Restart after a multi-day hiatus · Ken Tombs with Claude (design lead, Opus 4.8) · This note opens the Session 7 log (Entries 082–083) and recaps the pilot's aims through to the introduction of dimensions and the scorecard.

Prepared as the reconnection artefact after a genuine gap — itself an experimental condition, below.

Opening log entries

Entry 082 — Session 7 opened after a multi-day hiatus — recorded as an experimental condition

The pilot resumes after a genuine gap of several days — the first session boundary longer than a few hours (prior boundaries were overnight or a lunch break). The hiatus is logged not as a scheduling fact but as a condition: it simulates the natural interruptions of real litigation — evidence preparation halted mid-matter, for example by a judge calling the parties to chambers, or a matter parked pending a ruling, or counsel pulled onto something urgent.

What it tests, that no prior session did: whether the custody architecture reconstructs a verified working state from RECORDS ALONE after a real gap, when neither operator nor design lead retains the details in working memory. This is the direct, live test of the pilot's founding doctrine — "the chat is a convenience; the document is custody." Until now that doctrine was asserted; the hiatus makes the wake-up ritual prove it. The freeze verification gains a second dimension too: not merely "did the configuration hold overnight" but "did a frozen subject stay frozen across a genuine hiatus" — the first real evidence on whether an evidence-team AI can be relied upon to be the same colleague after a fortnight away.

Entry 083 — Design-lead model confirmed: Opus 4.8 through Session 9, Fable 5 from Session 10

The design lead (the Claude instance directing method, drafting, forensics, and logging — not the experimental subject) runs on Opus 4.8 for Sessions 7–9, on cost-capability grounds: procedure design, forensic reasoning over pasted outputs, and document generation do not require frontier capability, and Opus economises the resource pool that heavy subject sessions also draw on — the pilot's own break-even logic applied to its own staffing.

Escalation to Fable 5 is reserved for Session 10 (the multi-dimensional primer / axiom / scorecard analysis — the densest reasoning of the arc, layering primers and scoring against independent logs) and for any implementation stretch (e.g. Session 8's server build) that demonstrably needs it, as a deliberate, logged act.

Doctrine preserved without exception: the design-lead model choice NEVER touches the experimental subject. Claude Code remains on Sonnet default throughout, per Entries 009/040 and the frozen-subject configuration deployed in Session 6. "Out of the box" means the unconfigured default; the subject is now literally pinned by settings.json. Whatever model the design lead runs, the subject is Sonnet.

Recap — the pilot's aims through to dimensions and the scorecard

The pilot asks a single question in stages: can agentic AI be trusted with legal evidence, and what makes it trustworthy? Sessions 1–9 answer the infrastructure half — capability, governance, custody, and mediation. Session 10 opens the method half — the primer apparatus itself. The table records each stage's aim and where it stands at this restart.

Session	Aim	Status / outcome
1–5 (complete)	Establish and close the comparison design: can out-of-the-box AI do evidence work, and what does a primer add?	DONE. Retrieval verified 9/9 precision and recall against independent ground truth under three conditions (primed/4.6, baseline/5, primed/5). Capability constant; the primer's contribution isolated as exhibit discipline — accurate work becomes citable work. Governance overhead measured ($\approx 3\times$ approvals under the primer). 70 entries.
6 (complete)	Harden the boundaries: wake-up and sleep as documented procedures; recover what was lost; freeze the subject.	DONE. Four procedures drafted and two executed; Session 5 fully harvested; the hidden .git solved (custody instrument broken by the rename it might have recorded); the destroyed report recovered from the tool's checkpoint store; the frozen-subject configuration deployed — and the volatility's root cause found in a settings file ("autoUpdatesChannel": "latest"). 11 entries (071–081).
7 (opening)	Design the laboratory bench on paper: ADR-004 — the MCP custody server, frozen before any code.	Tool set (read_file, list_folder, search_messages, get_headers, hash_document, verify_hash); self-describing definitions; scoping and privilege-gate rules; chained custody-log schema. Side-deliverable: the Session 10 scorecard rubric, fixed in advance.
8	Build and smoke-test the server; connect Claude Code in passthrough mode (Phase 2a).	Working server; first mediated run; plumbing proven. Design-lead escalation to Fable 5 available if the code strains.
9	Measured Phase 2 runs: 2a passthrough vs 2b custody-on, against the existing grid.	ADR-003's central number — the cost of governance, isolated. The trace-versus-log gap measured for the first time (an infrastructure log finally exists).
10 (the turn)	The multi-dimensional analysis: multiple primers in concert, axioms as a distinct stratum, and the scorecard applied.	Where governance-as-method meets governance-as-infrastructure. Design lead escalates to Fable 5. Beyond this point the pilot measures the primer apparatus itself, scored against Session 9's independent logs.

The through-line, in one paragraph

The pilot has established that a stock AI retrieves evidence with verified accuracy, that a primer converts accurate work into citable work (exhibit discipline), and that the real risks live not in capability but in accountability — memory that persists invisibly, versions that drift silently, authority that blurs, and records that must be built rather than assumed. Sessions 7–9 answer these with infrastructure: a mediation layer (the MCP custody server) that enforces policy structurally and generates a custody record independent of anything the AI reports about itself, measured against the Phase 1 baseline. Session 10 then turns from infrastructure to method — introducing multiple primers in concert, axioms as a distinct governing stratum, and a scorecard that makes "exhibit discipline" a measured property scored against Session 9's independent logs rather than a characterisation. Everything up to Session 10 builds the bench; Session 10 begins to measure what the analyst does at it.

Immediate agenda for Session 7

First, the wake-up ritual — now carrying real experimental weight, as the first reconstruction of state from records across a genuine hiatus. Within it, the freeze verification: launch once, read the banner, and learn whether the frozen-subject configuration held (expected: still 2.1.209 / Sonnet 5; no auto memory; no prompt-suggestion ghost text). Then brief housekeeping (deduplicate GIT Deliverables;

locate the mislaid break-even note; reissue the procedure v1.0s with Session 6's earned amendments). Then the session's substance: ADR-004, the custody server designed and frozen on paper, with the scorecard rubric drafted alongside it.