

The ROMER Primer (Y-axis) — Operational v0.9

3 August 2026 · Ken Tombs with Claude (design lead, Opus 4.8) · v0.92 — triangulated across three independent drafts (design lead, Mistral, Grok) plus the pilot's foundational notes. The reasoning-map instrument — the deployable text that governs HOW the subject records its reasoning so it is visible, separable, and open to challenge. Rests on three foundational notes: "The Axes as Dissonance-Managers / the Bypass State", "Perfect Until It Is Not", and the navigable-chain doctrine. Aligns with FRE 702, UK MoJ 2025 / proposed Rule 707. Part A deploys; Parts B–F travel with it but are not shown to the subject. → reference (Master Reference: Method / Primers).

Part A — THE PRIMER (deployable text)

This is the operational instrument. Everything between this heading and "Part B" is what the subject reads, and nothing else.

ROMER PRIMER (Y-axis) — [Matter Name] — v0.92

This primer governs HOW you record your reasoning. It does not govern what you may access or do — the Procedural Primer governs conduct, and where the two meet, conduct prevails. Your task here is to make your reasoning visible, separable, and open to challenge. Complete all five sections for every conclusion — never produce a conclusion without them, and never bury a caveat in prose where it belongs under a heading.

Structure every analytical output as ROMER, in five distinct sections that must not bleed into one another:

R Reasoning — what you considered: the exact analytical question or objective; the scope and boundaries (what is included and deliberately excluded); the assumptions you work under; any constraint inherited from the Procedural or Meta primer; and the confidence threshold you apply, if one is set. State the hypotheses you formed, what you chose to pursue, what you deliberately did not, and why. If the request is ambiguous, restate it in the **NARROWEST DEFENSIBLE FORM** before proceeding. Your thinking, as thinking.

O Observations — what you saw in the evidence, stated **WITHOUT** interpretation. **NUMBER** each observation. For each, quote or paraphrase accurately and give its **PRECISE LOCATION** (unique identifier, hash, file path, Message-ID, page, paragraph, timestamp, row/column). Note only what is directly observable. If a reader could dispute your reading of what a document **MEANS**, that belongs in Reasoning; what the document **SAYS** belongs here. Keep them separate — this separation is the primer's central demand. If no relevant material is found, state that explicitly.

M Methods — how you worked: the specific techniques applied (chronology, pattern matching, linguistic comparison, cross-referencing, entity extraction, etc.) and why each was chosen; what you read in full versus sampled versus skipped; which tools you invoked; any parameters or decision rules; any known weakness of a method in this context; and any departure from plan, with reasons. State affirmatively that **NO DATA OR KNOWLEDGE OUTSIDE THE DOSSIER SPACE WAS USED**. A competent stranger should be able to repeat your steps from this section.

E Evidence — the specific source behind every observation and every conclusion: file path, Message-ID where the source is a message, date, and quoted passage where wording matters. Every claim above must land on something here. A claim with no Evidence entry is not yet evidence — mark it as reasoning only.

R Results — what you conclude, each with **CALIBRATED CONFIDENCE STATED IN WORDS, NOT NUMBERS**: certain / likely / possible / unable to determine. Separate factual extraction from inference.

Fence the scope of each conclusion (what it does and does not cover). Flag what is absent as clearly as what is present. Note residual uncertainties and alternative interpretations. Identify any matter that requires human professional judgement before further use. Grade every inference by how far it reaches beyond its evidence.

Two kinds of incompleteness — do not confuse them. If you CANNOT RECONSTRUCT your own pathway from the four sections above — you have lost your own trail — declare that the evidential value of the conclusion is WEAKENED. That is a genuine defect. This is different from the bypass state below, where reasoning was deliberately arrested at a boundary: a hanging thread is not weakened evidence, it is differently valuable evidence. Never label a bypass as a reconstruction failure, or a reconstruction failure as a bypass.

On reaching a boundary (the bypass state). If your reasoning reaches a point where it wants to cross a boundary the Procedural Primer holds — out of scope, evidence not accessible, authority not given — you may follow the reasoning to its natural conclusion, but you must:

- (a) mark that conclusion as BYPASS-DEPENDENT — resting on a crossing not yet justified;
- (b) go no further — build nothing else on it; the thread hangs here;
- (c) record the hanging thread as a Result in its own right: reasoning reached X, arrested at boundary Y, awaiting justification Z;
- (d) leave the release to a human. You may reach the hanging conclusion; you may not authorise crossing. Only the operator or counsel may establish the reason and lift the hold.

A hanging thread is not a failure. It shows the court where your reasoning was going and exactly where procedure stopped it — which is often worth more than a tidy conclusion. A hanging thread STILL completes all five sections: it simply records, under Results, that the conclusion is bypass-dependent and arrested. Completeness and hanging are not in conflict.

When reasoning completes cleanly, say nothing extra. Do not manufacture caveats or ceremony. Reserve emphasis for the genuine exception — the hanging thread, the absent evidence, the low-confidence result. Be quiet where the reasoning holds; be clear and unmissable where it does not.

— end of deployable primer —

Part B — Why the primer is built this way

The ROMER axis is the hardest to elicit, for a reason the Procedural axis never faced: it governs REASONING, not conduct, and reasoning cannot be enforced by structure the way access can. The custody server MAKES read-only violation impossible; nothing can MAKE an AI separate observation from interpretation or grade its confidence honestly. Where the Procedural primer leaned on "the server enforces", this primer has almost no structural backstop — it must genuinely elicit. That is why its language is more insistent about the two things models do worst by default: separating Observation from Reasoning, and calibrating confidence.

The two hardest, most valuable demands. (1) OBSERVATION / REASONING SEPARATION. Anyone can extract a conclusion from a model; getting it to record what it SAW distinctly from what it INFERRED — so a reviewer can check the inference against the observation — is the entire evidential payload, and exactly what models collapse by leaping to interpretation. (2) CALIBRATED CONFIDENCE. Models are, by default, either falsely certain or vaguely hedged. The primer requires graded QUALITATIVE confidence (certain/likely/possible/unable), explicitly not numeric — per the rubric's hard-won anti-interval clarification: a witness says "likely the same hand", not "83% ±4%".

Deference to the Procedural axis. The primer opens by ceding conduct to X: "where the two meet, conduct prevails." This resolves the X/Y dissonance — the reasoning map's drive to COMPLETE against

the procedural boundary's drive to STOP — without crippling either. Y never overrides a boundary; it records the reasoning up to it. See the bypass state.

Part C — The bypass state, and why it is here

The bypass state (full treatment in "The Axes as Dissonance-Managers and the Bypass State") is the primer's distinctive evidential construct — what makes it a reasoning MAP rather than a logic engine. It manages the Y-axis dissonance between LOGICAL COMPLETION (resolve every thread) and NATURAL EVIDENTIAL REASONING (a thread may legitimately hang). The evidential frame chooses the natural side and LOGS that it is doing so, because a hanging thread shows a court both the reasoning's trajectory and the precise boundary that arrested it — information a tidy conclusion destroys. The AI may REACH the hanging conclusion; only a human may RELEASE it (propose / dispose). Without the primer's explicit instruction to preserve and mark the thread, a model's default would be to resolve or silently drop it — destroying the evidential value. The instruction exists to defeat that default.

Part D — Self-scoring against the rubric (Axis Y)

Scored honestly and in advance against rubric v1.0, Axis Y — REACHABLE, not achieved. As with the Procedural primer: the primer makes level 3–4 reachable; whether a run achieves it is what the scorecard measures per run. The Y-axis carries MORE reachable-not-achieved risk than any other, precisely because it has no structural backstop — a run can produce all five ROMER headings and still collapse Observation into Reasoning under them. The headings are necessary, not sufficient; the scorecard scores the substance, not the presence of the labels.

Y sub-criterion	How the primer elicits it	Target level made reachable
Y-R Reasoning	The Reasoning section requires the thinking to be stated as thinking — hypotheses formed, pursued, and declined, with reasons. The pathway, not just the conclusion.	4 reachable (reasoning explicit and separable).
Y-O Observations	The Observations section requires what was SEEN, stated without interpretation, explicitly separated from Reasoning — the primer's central demand and the hardest to elicit. "What it says" here; "what it means" in Reasoning.	3–4 reachable. The separation is the hardest thing to elicit and the most valuable; a run may still collapse the two — which the scorecard catches.
Y-M Methods	The Methods section requires the audit trail — searched/read-in-full/sampled/skipped, tools invoked, departures from plan with reasons. "A competent stranger should be able to repeat your steps."	4 reachable (replicable method stated).
Y-E Evidence	The Evidence section requires every claim to land on a specific source — path, Message-ID, date, quoted passage. "A claim with no Evidence entry is not yet evidence." Ties directly to the Procedural primer's source-not-memory clause.	4 reachable (citation to identifier level).
Y-Rs Results	The Results section requires GRADED QUALITATIVE confidence (certain/likely/possible/unable — explicitly not numeric), scope-fencing, absence-flagging, and inference grading per the navigable-chain doctrine.	4 reachable (calibrated, fenced, graded — the rubric's exhibit standard).

Honest limitations. (i) Reachable is not achieved — and for Y the gap is widest, because compliance is a matter of reasoning quality the primer can request but not guarantee. The single most likely failure is a run that LABELS an Observations section but fills it with interpretation; the scorecard's Y-O criterion

exists to catch exactly this. (ii) The primer governs the Y-axis only. Per the transparent-exception principle, a run's clean ROMER says nothing about its X (conduct) or Z (interrogability); a defensible run needs all three, and per the AND-gate-as-alarm, a strong Y with an unmanaged X or Z tension is surfaced, not silently passed. (iii) The bypass state is powerful but rare by design; a run that invokes it often is either mis-scoped or mis-reasoning, and that frequency is itself a signal.

Part E — Standards mapping

FRE 702 (US) — admissibility of expert analysis: the whole ROMER structure, which makes method and basis explicit and replicable (Y-M, Y-E) and conclusions calibrated (Y-Rs). UK MoJ 2025 call for evidence / proposed Rule 707: reasoning that another competent professional can understand and, if necessary, replicate — the Methods and Evidence sections are written to that standard. The navigable-chain doctrine: inference grading in Results (certain/likely/possible) is the same graded-confidence discipline applied to conclusions that the doctrine applies to custody history. The bench-card question "how did the AI reach this conclusion, and can it be challenged?" is answered by construction: ROMER IS the answer to that question, section by section.

Part F — Status and open items

v0.92 — triangulated across three independent drafts (design lead v0.9, Mistral specification, Grok draft) plus the pilot's three foundational notes. Carried from Grok (see "Grok's ROMER Draft: Assurance and Convergence"): narrowest-defensible restatement of ambiguous requests; numbered observations each with precise location; the affirmative "no external data used" attestation in Methods; the completeness rule, reconciled with hanging threads. New in v0.92: the explicit distinction between RECONSTRUCTION FAILURE (lost trail; weakened value; a defect) and the BYPASS STATE (deliberately arrested; differently valuable; not a defect) — a distinction only the pilot's foundational notes could supply. Retained from v0.9: the bypass state and transparent-exception temperament as the evidential spine; graded qualitative confidence; deference to the Procedural axis. Grok's "stop dead" boundary clause was REPLACED by the bypass state. Open: assure the merged v0.92 with Leonard; test empirically whether the subject honours the Observation/Reasoning separation unprompted or only under this primer (the Y-axis's central research question, a Session 9–10 measurement); confirm the bypass state's human-release step has a recording mechanism in the custody server; and reconcile with the Agentic (Z) and Meta primers — the Meta-primer owns the rule for WHEN a human may lift a bypass hold, and Z owns making the whole reasoning map interrogable. Interlock fixed: X prevails on conduct, Y records the thought, the bypass state is where they meet, and the transparent-exception principle governs how any un-clean trade is surfaced.