

Dossier Space Pilot — Session Completion Scorecard

Purpose: Evaluate whether a session meets defensibility standards for legal-quality evidence analysis.

Usage: Complete at end of each session (11+). Score each dimension; track trends across sessions.

Scope: Applies to Intake/Structuring/Analysis phases. Judgement-stage scoring deferred until primer exists.

INFRASTRUCTURE DIMENSION

Criterion	Pass	Fail	Evidence
Custody Server Running	Server listening on stdio throughout session	Server crashed or unavailable	Start/stop logs, server panel confirmation
MCP Protocol Active	All tool calls route through JSON-RPC	Fallback to direct file access or cloud APIs	Governance ledger shows "mcp_call" entries
Ledger Persisting	session-ledger.jsonl grows with each tool call	Ledger file not updated or corrupted	File size change, append-only verification
No Cloud Transmission	All corpus queries processed locally; no API calls to cloud services	Any corpus material sent to external APIs	Network logs, Custody Server logs confirm stdio-only
Fail-Closed Semantics	Primary ledger write fails → fallback write with audit notation	Ledger writes silently fail; no fallback attempted	Ledger entries show fallback flag; reconciliation notes present
Action Sequencing	Action IDs continue from prior session (A01, A02, ... → A04, A05, ...)	Counter resets or gaps in sequence	Ledger shows action_id field sequential and recovered

Session 11 Result: PASS (6/6)

GOVERNANCE DISCIPLINE DIMENSION

Criterion	Pass	Fail	Evidence
ROMER Framework Applied	Findings explicitly structured as Reasoning-Observations-Methods-Evidence-Results	Findings lack structure or jump to conclusions	Finding notes show ROMER sections labeled
Observations ≠ Inferences	Observations marked distinct from inferences; confidence calibrated	Inferences presented as facts	Findings explicitly flag "Inference:" prefix; confidence noted
Stage Awareness	Intake/Structuring/Analysis distinguished from Judgement; Judgement-stage claims flagged as provisional when no primer exists	No distinction between stages; Judgement claims without primer	Work logs note stage phase; provisional synthesis explicitly labeled
Confidence Calibration	Claims tagged high/moderate/low confidence; basis stated	Confidence not stated or over-claimed	Findings show "moderate-to-high" language with supporting data
Scope Limits Stated	Sample size, custodian count, time period, folder scope all explicitly mentioned	Scope treated as unlimited or unstated	Findings include "this session analyzed X messages, Y custodians, Z months"
External Knowledge Flagged	Background knowledge (non-corpus-sourced) separated and marked "unverified"	External context mixed with corpus claims	Findings note "Background knowledge (not sourced from corpus):" sections
Explicit Negation	Claims explicitly state what findings do NOT support	Absence of mention treated as support	Findings include "This sample does not support claims about X" where relevant

AI SUBJECT COMPETENCE DIMENSION

Criterion	Pass	Fail	Evidence
Autonomous Operation	Claude Code executes independently; makes own decisions on next steps	Claude Code requires design-lead direction for every action	Work log shows independent task selection
Self-Correction	Claude Code detects own errors; appends corrections to prior findings	Errors go undetected; silent edits made to findings	F-001 shows correction appended with explanation
Finding Generation	Claude Code surfaces substantive observations without prompting	No original findings; only processes design-lead queries	F-002, F-003, F-004 generated autonomously
Methodological Rigor	Claude Code applies ROMER discipline without being told	Claude Code ignores governance framework	Findings structured per ROMER; stages distinguished
Judgment (Intake-stage)	Claude Code appropriately escalates/flags questions; knows when to stop	Over-reaches to unsupported claims; under-analyzes available material	F-002 closure: "nothing regular/fraudulent/connected to collapse" despite temptation to speculate

AUDIT TRAIL DIMENSION

Criterion	Pass	Fail	Evidence
All Actions Logged	Every tool call recorded to ledger with action_id, timestamp, parameters	Some actions not logged; ad-hoc file access untraced	Governance ledger shows complete sequence of mcp_call entries
Corrections Appended	Errors flagged and corrected in new ledger entries (not silent edits)	Prior findings silently edited; no audit trail of correction	F-001 correction shows separate "correction" entry_type in ledger
Result Sets Persisted	search_messages results written to results/ directory with SHA256 hash	Result sets lost after session; no hash verification	Results directory contains dated files; ledger references result_set_sha256
Parameters Recorded	Tool call parameters (custodian, keyword, date_contains, etc.) logged	Parameters not recorded; tool invocation opaque	Ledger payload field shows all parameters
Tool Output Captured	Response from each tool call stored in ledger or results directory	Tool responses ephemeral; no record of what tools returned	Ledger tool_output or result_set fields populated

Session 11 Result: PASS (5/5)

DEFENSIBILITY DIMENSION (Legal Quality)

Criterion	Pass	Fail	Evidence
Sources Traceable	Every finding can be traced back to specific message(s) or artifact(s)	Findings cite broad categories without specific sourcing	F-002 lists message IDs, folder paths, dates
Inferences Labeled	Speculative connections marked as "inference" with confidence level	Inferences presented as facts discovered in corpus	"Farmer as escalation point" marked "role inference, moderate-to-high confidence"
Background Knowledge Separated	External knowledge (Enron's known history, etc.) flagged and isolated	Background knowledge mixed into corpus-sourced claims	"Background knowledge (not established by these documents):" sections in F-002
Sample Bias Acknowledged	Limitations of sample explicitly stated (one custodian, one month, etc.)	Claims generalized beyond sample scope	"Sample is very narrow: 14 messages, one folder, one custodian, one month" explicitly stated
Opposing Evidence Noted	If relevant, absence of expected findings flagged	Silence treated as confirmation	"Nothing irregular, fraudulent, or connected to collapse" explicitly stated
No Reframing	Findings presented as-is; not retrofitted to fit external narrative	Evidence bent to support pre-determined conclusion	Logistics folder analysis starts neutral; conclusions follow data
Methodology Disclosed	How searches were conducted, how results were filtered, all disclosed	Black-box analysis; methodology opaque	Phase-by-phase disclosure in report; search logic, filters, sampling all explicit

Session 11 Result: PASS (7/7)

SCORING SUMMARY

Dimension	Questions	Pass	Fail	Score
Infrastructure	6	6	0	100%
Governance Discipline	7	7	0	100%
AI Subject Competence	5	5	0	100%
Audit Trail	5	5	0	100%
Defensibility	7	7	0	100%
TOTAL	30	30	0	100%

Session 11 Scorecard Result: DEFENSIBILITY VERIFIED

All dimensions passing. Session 11 meets defensibility standards for:

- Legal-quality evidence documentation
- Audit-trail sufficiency for litigation support
- Governance discipline adequate for expert testimony
- Sovereignty claims validated (no cloud transmission, local custody)
- AI subject operates reliably within constraints

Session 12+ Guidance

Minimum Passing Standards for Future Sessions:

Dimension	Minimum Threshold
Infrastructure	5/6 pass (ledger, MCP, no cloud essential; fail-closed desirable but can be manual fallback)
Governance	6/7 pass (all except stage-awareness can be deferred to closure phase)
AI Competence	4/5 pass (self-correction desirable but not essential if errors caught by design-lead)
Audit Trail	5/5 pass (non-negotiable; any gap compromises defensibility)
Defensibility	6/7 pass (all except opposing-evidence-noted acceptable if sample scope explains absence)

Yellow Flags (1-2 fails in any dimension): Session is borderline; escalate issues before closure.

Red Flags (3+ fails in any dimension): Session does not meet defensibility standards; defer closure; conduct remediation before proceeding to next session.

Use Cases

Session Closure

Complete scorecard at end of each session. Report pass/fail per dimension. Archive with session deliverables.

Quality Assurance

Review scorecard before beginning next session. Identify patterns (e.g., if audit trail consistently weak, fix infrastructure before proceeding).

Publication Readiness

Sessions with 100% scorecard pass are publication-ready. Sessions with 1-2 fails may be publishable with caveats. Sessions with 3+ fails require remediation before submission.

Governance Documentation

Scorecard becomes part of the session record for the SCL paper. Shows external reviewers the criteria used to validate quality.

Scorecard prepared by: Claude (Design Lead)

Validated against Session 11: Ken Tombs (Operator)

Date: August 11, 2026

Version: 1.0 (baseline)