

## Session Planner — v1.0

15 July 2026 · Ken Tombs with Claude · Standing plan: amendable per session; every amendment produces a new version of this document (v1.1, v1.2 ...), forward-only.

### 1. Where the pilot stands

Sessions 1–5 (13–14 July) completed the comparison design: the same retrieval task verified at 9/9 precision and recall under three conditions — primed on Sonnet 4.6, baseline on Sonnet 5, primed on Sonnet 5 — isolating the primer’s contribution within one model generation. Central claim demonstrated: the primer restores the frame of exhibit discipline; capability is constant; governance overhead is measurable (approvals  $\approx 3\times$  under the primer). Seventy log entries, seven documents, zero subject writes into the corpus. Full detail: the five session logs and the Pilot Synthesis.

### 2. Sessions completed

| # | Date           | What happened                                                                                                                                                                                                                                                                                                                                                                                                 | Deliverables                                      |
|---|----------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|---------------------------------------------------|
| 1 | 13 Jul pm      | Ground rules; workspace scoping (one contamination event contained); primer created (verification later found to be theatre — v0.1 governed); baseline recorded (2.1.140 / Sonnet 4.6); Run 1.1 survey delivered; filename anomaly tested — resolved in subject’s favour; rename provenance question opened. Entries 001–018.                                                                                 | Session 1 log; ADR-003 (drafted earlier same day) |
| 2 | 13 Jul eve     | Run 1.2-W retrieval 9/9 verified against independent ground truth; cold start failed (invisible auto-resume); subject found actively excavating its history; four unrequested memory files written (incl. operator profile) — read, archived, deleted; v0.1 discovered in force all day — runs relabelled primed; rename attributed, per-file log found empty. Entries 019–033.                               | Session 2 log; memory exhibit                     |
| 3 | 14 Jul am      | Wake-up ritual born; transcript quarantine (one malformed copy caught); primer swap byte-verified; subject mutated overnight (2.1.207 / Sonnet 5, 4.6 deprecated) — design rebuilt on “Claude Code current”; Run 1.1-B delivered then VOIDED: transcripts found in workspace — forensics exonerated subject, attributed the audit’s own truncated copy command; hidden .git discovered. Entries 034–045.      | Session 3 log; transcript exhibit                 |
| 4 | 14 Jul late am | True baseline established: Run 1.1-B2 survey verified cold; Run 1.2-B retrieval 9/9 (work preceded the peek); tool’s own governance caught a /tmp write before the primer was tested; filename collision destroyed the 13 Jul report (recovery pending); primed-vs-baseline delta characterised — accurate vs citable; suggestion feature found drafting prompts into the authority channel. Entries 046–053. | Session 4 log; run 1.2-B/W outputs exhibit        |

| # | Date      | What happened                                                                                                                                                                                                                                                                                                                                                                                                                                                                                | Deliverables                                                                |
|---|-----------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------|
| 5 | 14 Jul pm | Primed re-run on Sonnet 5 — comparison closed within one generation. En route: the bench that wasn't (routing failures explained); wrong-floor launch (no primer loads — floor determines behaviour); --session-id fails silently; tool storage anatomy moved (sessions/file-history/backups/paste-cache). Runs 1.1-P5 and 1.2-P5: dual-source verification, ROMER by name unprompted, 9/9, deleted_items cluster discovered. Approval overhead $\approx 3\times$ baseline. Entries 054–070. | Session 5 log; Pilot Synthesis; break-even reference note; two P5 artefacts |

### 3. Sessions planned

| #  | Title                              | Scope                                                                                                                                                                                                                                                                                                                                                                                               | Deliverable                                    | Credit                         |
|----|------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------|--------------------------------|
| 6  | <b>Boundaries (domestics)</b>      | Wake-up and sleep procedures executed end-to-end, timed, and documented; cold/warm start verification; frozen-subject investigation (can auto-update, suggestions, memory be disabled?); Session 5 transcript harvest; zero-credit recoveries — .git inspection (possible per-file rename custody record) and file-history recovery of the overwritten 13 Jul report.                               | Wake-Up & Sleep Procedures v1.0                | Near zero                      |
| 7  | <b>ADR-004: the bench on paper</b> | MCP custody server designed before any code: tool set frozen (read_file, list_folder, search_messages, get_headers, hash_document, verify_hash); self-describing tool definitions per Leonard's spec; scoping and privilege-gate rules; chained custody-log schema. Side-deliverable: draft the Session 10 scorecard rubric (criteria fixed before the runs they will score).                       | ADR-004 + frozen tool spec; scorecard draft    | Minimal                        |
| 8  | <b>Build &amp; smoke-test</b>      | Dossier Space MCP Server v1 implemented in Python to the frozen spec (Claude writes, operator deploys); Claude Code connected in passthrough mode (Phase 2a, custody off); one test task through the pipe to prove the plumbing.                                                                                                                                                                    | Working server; first mediated run             | Moderate                       |
| 9  | <b>Measured Phase 2 runs</b>       | Survey and retrieval through 2a (passthrough) and 2b (custody on), scored against existing ground truth and the Phase 1 grid. ADR-003's central number measured: the cost of governance, isolated. Trace-versus-log gap metric runs for the first time (an infrastructure log finally exists).                                                                                                      | Completed experimental grid, all phases        | The credit session — budget it |
| 10 | <b>The primer apparatus</b>        | Multiple primers in concert: a CIPDA-Intake stage primer drafted and deployed alongside standing v0.1 (the §4.2 layering, exercised for the first time); axioms as a distinct stratum (standing invariants vs stage procedure — where does each rule belong?); the scorecard applied — exhibit discipline scored per run against the pre-fixed rubric, so "the primer restores the frame" becomes a | Scored primer-layering runs; rubric v1 results | Moderate                       |

| #  | Title                                         | Scope                                                                                                                                                                                                                                                | Deliverable                | Credit |
|----|-----------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|----------------------------|--------|
|    |                                               | measured claim. Primer effects checked against Session 9's independent infrastructure log.                                                                                                                                                           |                            |        |
| 11 | <b>Synthesis v2 &amp; the turn to writing</b> | Grid and Synthesis updated across all phases and primer conditions; deleted_items cluster sampled if budget allows (~20 files, not 1,122); remaining open items closed or formally carried; SCL paper skeleton opened on the complete evidence base. | Synthesis v2; SCL skeleton | Cheap  |

#### 4. The logic of the ordering

Session 6 hardens the boundaries, where every serious incident of the pilot occurred. Session 7 designs on paper while credit rests — reversible work first; nothing is built on an unfrozen spec. Sessions 8–9 spend the credit where it buys the paper's missing measurement: the cost of governance, isolated, and the trace-versus-log gap. Session 10 turns from governance-as-infrastructure to governance-as-method, with the scorecard fixed in advance and Session 9's independent log available to check primer effects against. Session 11 banks everything and opens the writing. Standing caveats: the subject will likely update several times across this arc — if Session 6's frozen-subject investigation succeeds, Sessions 8–9 run on a pinned version; if not, "Claude Code current" doctrine continues and drift is recorded as environment. Credit is sequenced, not interleaved: heavy subject sessions and heavy drafting draw on the same pool.

#### 5. Standing doctrine (pointer)

The working rules — session/run/entry clocks, wake-up ritual, two-window discipline, verification outside the subject, per-command approvals, archive-then-delete-then-verify, per-run filenames, filed-means-frozen, identity verified at birth, corrections forward-only — are consolidated in the Pilot Synthesis §3 and are not restated here. This planner tracks WHERE the pilot is going; the Synthesis records HOW it travels.