

Good — let's set the ground rules before anything gets touched. Here's what I propose, and you amend as you see fit.

Roles. One thing to be clear about first: I'm Claude in the chat app, so I can't reach your workstation directly. Claude Code on your desktop is the experimental subject and does the touching; you are its operator; I lead the design, write what needs writing, and tell you exactly what to do at each step. That division is actually correct for the experiment — the subject stays out of the box, and nothing I do here contaminates it.

One step at a time. I give you a single instruction, in plain language, and wait. You do it, tell me what happened — success, error message, or something odd — and only then do we move on. No batches of commands, no assumed steps. If anything surprises you, we stop and record it before fixing it.

Everything is logged. I'll maintain a running Pilot Log — dated entries recording what we did, why, what happened, and any friction encountered. Friction is data, per your method, so we record failures as carefully as successes. At the end of each working session I'll produce the log as a document you save into the pilot folder. Between sessions I can search our past conversations to reconstruct where we left off, but the document on your disk is the authoritative record — my memory is a convenience, not custody.

Nothing changes silently. Any decision that alters the design — tool set, phase order, task definitions — gets recorded, and if it's significant, it becomes an ADR amendment rather than a quiet drift. The corpus is never modified; if we need working copies of anything, they go in a separate folder, never in the evidence tree.

Versions recorded before we start. First practical act of the experiment will be recording the Claude Code version, model version, Windows build, and the corpus location and state. That's the baseline entry in the log.

Claude Code stays stock. No custom instructions that enhance its capability during measured runs. The CLAUDE.md question needs a decision from you here, actually: a CLAUDE.md in the pilot folder is arguably a primer, which belongs to Phase 3a, not Phase 1. So for the Phase 1 baseline I'd suggest either no CLAUDE.md at all, or a minimal one containing only safety constraints (read-only, stay in this folder) — and we record which choice we made and why. This is the first real design decision of the run.

Do those rules suit you? And which way do you want to go on CLAUDE.md for Phase 1 — none, or safety-only?