

The Axes as Dissonance-Managers, and the Bypass State

3 August 2026 · Ken Tombs (whose insight this note records) with Claude (design lead, Opus 4.8) · A foundational reframing of what the three axes ARE, surfaced during the design of the ROMER (Y) primer, and the specific construct — the bypass state — it produces. This note underpins the ROMER primer and reframes the role of the Agentic (Z) primer. → reference (Master Reference: Architecture / Foundational concepts).

1. The reframing: the axes manage dissonance

Until now the three axes of Dossier Space have been described as three layers of evidential governance — procedural, reasoning, agentic. This note records a deeper account of what they are, exposed for the first time during the ROMER primer design: the axes are not three CHECKLISTS that each verify a different kind of correctness. Each axis exists because some pair of genuinely contradictory things must be held together without collapsing one into the other. The axes are MANAGERS OF DISSONANCE.

Axis	The dissonance it manages	Why neither side may simply win
X — Procedural	Capability vs constraint — what the AI is ABLE to do against what it is ALLOWED to do.	Too much constraint blinds the analysis; too little unleashes it. The axis holds the line without pretending the capability is not there.
Y — ROMER	Logical completion vs natural evidential reasoning — the drive to RESOLVE against the honesty of the hanging thread.	Forcing resolution destroys evidentially-valuable incompleteness; permitting endless open threads destroys usable conclusions. The axis preserves the interrupted inference without abandoning the drive to conclude.
Z — Agentic	The necessary IMPERFECTION of X and Y vs the need for accountability.	Because X and Y can never perfectly resolve their tensions, the RESIDUE must be made interrogable. Z does not add a third correctness-check; it makes the management of the other two's imperfection visible and challengeable.

2. The honest corollary: the criteria are necessarily imperfect

If the axes manage real, irreducible tensions, then the axial criteria — the scorecard's sub-criteria, the primers' clauses — will inevitably be IMPERFECT. This is not a defect to be engineered away by cleverer criteria; the tensions are real, and no rubric can fully resolve a genuine contradiction in advance. Admitting this is a strength, not a weakness: a framework that claimed its criteria PERFECTLY ensured correctness would be making a false claim that the first hostile reviewer would puncture. Dossier Space instead claims something more defensible — that it MANAGES irreducible tensions and makes the management interrogable.

This is why the Agentic (Z) primer exists, and it dignifies Z's role. Z is not a third correctness-check bolted onto two others. Z exists PRECISELY BECAUSE X and Y are necessarily imperfect — it is the layer that makes the residue they cannot resolve visible, examinable, and challengeable. In a hypothetical world where X and Y perfectly resolved their tensions, Z would be redundant. In the real world where they cannot, Z is what makes the imperfection DEFENSIBLE: the imperfection is shown and interrogable rather than hidden. This also explains why the scorecard's gate is conjunctive (AND, not sum): the axes measure three different irreducible tensions, so a zero on any means one whole tension has gone unmanaged — which cannot be compensated by another axis managing a different tension well.

The framework is therefore closer to a jurisprudence than a quality-check: it admits that evidential reasoning under real constraint is irreducibly a matter of managing competing goods, and holds that the honest response is not to pretend the tensions away but to structure their management so it can be interrogated.

3. The Y-axis dissonance in detail: logical vs natural

The dissonance the ROMER axis manages is between LOGICAL COMPLETION and NATURAL EVIDENTIAL REASONING. Logical completion wants every inference resolved — reach the conclusion, close the loop; an unresolved thread is a defect. Natural evidential reasoning — how a careful investigator actually works — accepts that a line of thought can legitimately HANG: "I can see where this points, but I am not permitted to go there yet, so I mark it and hold." To a logician the hanging thread is a failure. To a court it is often the most honest and useful thing in the record, because it shows both the reasoning's trajectory AND the precise boundary that arrested it.

The evidential frame consciously chooses the NATURAL side against the LOGICAL side — and must LOG that it is doing so, and why. We do not permit hanging reasoning silently; we record that we are deliberately tolerating a state a logician would call broken, because in evidence work the broken-looking state carries information the clean state would destroy. Without the recorded rationale, a downstream reviewer — human or AI — encountering a hanging thread may "fix" it: resolve it, delete it, tidy it — and destroy exactly the evidential value it was preserved for. The logged dissonance is what tells the next reader: this is deliberate, this is allowed, do NOT resolve it. (The same defensive-documentation instinct as the "reachable is not achieved" hardening of the Procedural primer, applied to reasoning rather than scoring.)

4. The construct: the bypass state

The Y-axis dissonance produces a distinct, named construct — the BYPASS STATE — for when reasoning reaches a boundary it wants to cross:

- (a) Reasoning that reaches a boundary may CONTINUE to its natural conclusion — but that conclusion is MARKED as bypass-dependent: explicitly flagged as resting on a crossing not yet justified.
- (b) No further reasoning may BUILD ON a bypass-dependent conclusion until the reason for the bypass is established and recorded. The thread is QUARANTINED, not deleted — it hangs, visible, at the point of interruption.
- (c) The bypass is a SPECIAL, RARE state — not a routine move. Its exceptionality is what makes it meaningful; if reasoning bypassed boundaries casually, the flag would be noise. It is reserved for where the reasoning genuinely has somewhere to go that procedure will not yet permit.
- (d) The hanging reasoning is RECORDED AS A RESULT in its own right — "reasoning reached X, arrested at boundary Y, pending justification Z" — because that unresolved state is evidentially valuable and must survive into the output, not be tidied away.
- (e) The AI PROPOSES; a human DISPOSES. The AI may REACH the hanging conclusion but cannot self-authorise crossing. Only a human (operator or counsel) may establish and record the bypass reason — a scope extension, a ruling, a documented decision — at which point the quarantine lifts and downstream reasoning may resume, citing the justification. Releasing a quarantined thread is a human act on the record.

5. How this resolves the X/Y interlock

The bypass state resolves the tension between the Procedural primer's refusal/stop clauses (X6–7) and the ROMER drive to complete, and it does so more elegantly than a flat "stop dead". X (procedural) still WINS ON CONDUCT: no further evidence is pursued, no downstream work proceeds, the boundary holds. But Y (reasoning) is not SILENCED: its trajectory is preserved up to and including the arrested conclusion. X stops the PURSUIT; Y records the THOUGHT. They occupy different layers — the boundary halts ACTION while the reasoning map preserves the interrupted INFERENCE — so they are not in conflict. This gives the Meta-primer a precise job: adjudicating WHEN a bypass-dependent thread may be un-quarantined — i.e., when the reason for the bypass has been sufficiently established (by a human) to let downstream reasoning resume. That is exactly the kind of inter-axis arbitration the Meta-primer exists for.

6. Why this matters for the pilot and the paper

This note records the moment the framework's self-understanding deepened: from "three governance layers" to "three managers of irreducible dissonance, whose necessary imperfection the agentic layer makes accountable." For the SCL paper this is a stronger and more honest thesis than correctness-assurance — it positions Dossier Space as a disciplined way of MANAGING the tensions inherent in AI-assisted evidence work, not a false promise of resolving them. The bypass state is the first concrete instrument of that philosophy: a mechanism that preserves evidentially-valuable incompleteness, logs the dissonance that justifies it, and hands the human the sole authority to release it — reasoning that hangs, honestly, where procedure arrested it.