The Primer Gap: Why We Build the Bridge

3 August 2026 · Ken Tombs with Claude (design lead, Opus 4.8) · A scan of the wider field, prompted by the operator's observation that nothing available aligns metrics to criteria — confirming the gap and fixing the design decision for the Procedural Primer. → reference (Master Reference: Method / Primers).

1. What prompted the search

Two independent drafts of a Procedural Primer — the operator's outline and an independent draft from Mistral — both came back as SPECIFICATIONS (measures, indicators, standards mappings) rather than as a deployable primer. The operator's diagnosis: this is not a drafting failure but a reflection of the field — nothing available anywhere aligns runnable metrics with the criteria they claim to measure. Rather than reason from inside the pilot's own bubble, the design lead scanned the wider world to test the claim.

2. What the wider world actually offers

The market splits into two non-overlapping camps, and neither is what evidence work needs:

What the market offers	What it actually is	Why it is not what evidence work needs
"AI prompts for lawyers / investigators" (Cellebrite, Tavrn, US Legal Support)	TASK prompts — summarise this deposition, map the elements of this claim. Governance = "a human validates the output afterwards".	No procedural discipline, no scope enforcement, no attestation. Chain of custody is explicitly DISCLAIMED as belonging to the underlying evidence, "not the AI layer above it". The AI's own conduct is ungoverned.
"AI governance frameworks / ISO 42001" (the standards ecosystem)	MANAGEMENT-SYSTEM standards — organisational policies, risk registers, lifecycle controls, audit checklists, Plan-Do-Check-Act.	Operates at "does your ORGANISATION govern AI responsibly", not "what does THIS AI do when it touches THIS evidence". A binder, not a runtime. Passes a desk review; says nothing about per-run conduct.
THE BRIDGE (what Dossier Space builds)	A deployable RUNTIME instrument the AI reads before working — short, imperative, machine-facing — that turns governance intent into per-run behaviour and attestations.	This is the missing middle. Nobody builds it, which is why no source aligns metrics to criteria. The pilot can, because it has the runtime (custody server), the version answer (frozen subject), and the hash answer (chained log).

3. The gap, named by the field itself

The most useful single sentence found, from an ISO/IEC 42001 practitioner guide: "ISO requires documented information and governance processes (the paperwork), and it expects operational control and monitoring (the runtime). If your AIMS is only a binder of policies, you'll pass a desk review and still get burned in production." The field has PAPERWORK (governance frameworks) and TASK-PROMPTS (commercial legal AI) and nothing that bridges them — nothing that converts governance intent into runtime behaviour the AI actually exhibits, per run, on this evidence. That missing bridge is exactly what the operator sensed and exactly what the Procedural Primer must be.

4. Two things the search gave us beyond confirmation

First — the courts are already asking our exact questions, which tells us what the primer's OUTPUT must answer. Admissibility sources converge: for AI-touched evidence, "chain-of-custody documentation of the prompt, model version, and output hash is necessary" — which is precisely the pilot's prompt-logging, model-versioning (the freeze), and hashing, named as the admissibility requirement. The US National Center for State Courts has issued JUDICIAL BENCH CARDS: targeted questions judges read from the bench about "the source, chain of custody, metadata, and verification", and a key distinction between "acknowledged" AI evidence and "unacknowledged" evidence potentially altered to mislead. Design consequence: the Procedural Primer's attestations should map directly onto the bench-card questions, so a run's output is pre-authenticated against what a judge will actually ask — something no framework in either camp does.

Second — prior art exists for recording the AI as a custody actor. A US patent (11,106,693 / 10,817,529) for digital-evidence-management chain-of-custody explicitly contemplates recording "artificial intelligence and/or machine learning algorithms, and/or non-human entities... which accessed a piece of digital evidence". This validates the custody-log design (the AI as a logged accessing entity is an established idea) and flags prior art for Leonard's attention in how the mechanism is described.

5. The decision

We build the BRIDGE, not another specification and not a task-prompt. The Procedural Primer will be a short, imperative, machine-facing RUNTIME instrument — deployable and byte-verifiable like CLAUDE.md v0.1, frozen per matter — governed by three principles drawn from the research:

(1) It is the runtime, not the binder. Deployable text the subject reads, not a 13-page essay. The rich specification content (indicator tables, ISO clause mappings, validation tests) is preserved as a SEPARATE Rationale & Standards-Mapping Note that travels with the primer for the desk-review / assurance audience — the pilot's established split of instrument-from-rationale (as with ADR-004 and the rubric).

(2) It elicits declaration; the server enforces. Per the pilot's core thesis (governance by structure, not vigilance): the primer makes the AI DECLARE and exercise discipline; the custody server makes violation IMPOSSIBLE. Read-only is not a request the AI might fail — it is absent from the toolset. The primer says "attest you accessed only through the custody tools", not "please do not write files".

(3) Its outputs answer the bench card. The required attestations — scope, integrity, provenance, constraints, refusals — are structured to be exactly what a judge or opposing counsel will ask, so the run's output is pre-authenticated against the questions courts are already publishing.

This is the artefact the field is missing and the pilot is uniquely positioned to build — because Dossier Space already has the runtime (the custody server), the version answer (the frozen subject), and the hash answer (the chained custody log). We are not writing a framework; we are writing the thing frameworks assume exists and never provide. Convergence of sources: the operator's outline, Mistral's draft, and the field scan all feed the two artefacts — the runtime primer and its rationale note — each source doing the job it is actually fit for; nothing wasted.