

Weighting the scorecard deployment – discussion with Leonard.

Ken, I think the four options are not mutually exclusive, but they belong at **different levels**. The mistake would be forcing one scorecard to do every job.

My view is that Dossier Space needs a **tiered scoring architecture**, with the question-level scorecard as the primary operational unit.

The four options weighed

1. One scorecard for the whole disclosure/discovery corpus

This would assess the overall Dossier Space or matter:

- Was the corpus admitted properly?
- Is provenance sufficiently understood?
- Are scope, custody and privilege controls functioning?
- Is the analytical environment generally reconstructible?
- Are the logs complete and tamper-evident?

Strengths

It gives a useful **matter-level assurance statement** and tells the court or legal team whether the environment as a whole is governed.

It is also efficient. The common procedural conditions—server version, corpus baseline, scope configuration, logging, read-only access—do not need to be rescored every time.

Weaknesses

It is too coarse to support an individual answer.

A corpus could be well governed while a particular answer is poorly reasoned. Equally, one defective item should not necessarily drag the entire matter score to zero.

The current rubric says scores must come from the filed record and custody log rather than impression. A whole-matter score could easily become a general certification rather than a test of what actually happened in a particular analysis.

Verdict

Necessary, but insufficient.

It should operate as a **Matter Assurance Baseline**, not as the only scorecard.

2. A scorecard for every question posed to Dossier Space

This is the strongest primary unit.

For example:

How many times did Person A communicate with Person B, and what was the essence of those communications?

That question generates a defined run or analytical package:

- declared intent;
- admitted scope;
- search method;
- items retrieved;
- exclusions and deduplication;
- ROMER analysis;
- resulting answer;
- custody trail;
- exceptions.

The scorecard can then ask:

- Was the question answered within scope?
- Was the method reproducible?
- Were claims tied to evidence?
- Were uncertainties surfaced?
- Can the result be reconstructed and challenged?

Strengths

It aligns directly with the actual unit of legal reliance: **the proposition or answer being used**.

It also fits the current design of the rubric, which speaks repeatedly in terms of scoring “each run.”

It allows different questions within the same matter to receive different scores. A simple communication count may be strongly defensible, while an interpretation of motive may remain weak.

It also captures coupling properly. Procedural controls may constrain one question without affecting another.

Weaknesses

A “question” can be very broad or very narrow. The system will need a rule for when a follow-up is:

- part of the same analytical run;
- a refinement;
- or a new scoreable question.

There is also some overhead, though much of X and Z can inherit from the matter baseline unless an exception occurs.

Verdict

This should be the primary scorecard unit.

Call it the **Question or Proposition Scorecard**.

3. A scorecard for every individual item of evidence

This would score every email, document, attachment or record against X, Y and Z.

At first this sounds rigorous, but I think it misapplies the architecture.

An individual email does not normally contain a ROMER reasoning pathway. It is evidence used in reasoning. Nor is every item independently “agentic” or interrogable in the same sense as an analytical output.

What can sensibly be assessed at item level

An evidence item can have an **evidence status record**:

- source;
- custodian;

- path;
- hash;
- date acquired;
- transformations;
- privilege status;
- provenance flags;
- authenticity limitations;
- relationship to duplicates or versions.

That is primarily procedural metadata.

Strengths

This is valuable where individual items are contested, transformed, incomplete or central to a conclusion.

It supports chain of custody and makes later challenge much easier.

Weaknesses

Applying the full X–Y–Z scorecard to every item would produce enormous volume and much false precision.

It would also blur the distinction between:

- the quality of an item;
- and the defensibility of the reasoning built from it.

A perfectly preserved email can still be misinterpreted. A provenance-limited item can still be useful if its limitation is properly exposed.

Verdict

Do not use the full scorecard for every item.

Instead, maintain a lighter **Evidence Item Status** or **Evidence Passport**. Escalate selected items to fuller assessment only where material.

4. Use the scorecard only when an item or answer is challenged

This is attractive because it follows the transparent-exception principle: stay quiet until something does not hold cleanly.

It would reduce cost and ceremony.

Strengths

It is proportionate and focuses attention where adversarial scrutiny actually occurs.

It avoids making every analytical act feel like a formal court exhibit.

It also fits the idea that the framework should be silent on the perfect and vivid on the flawed.

Weaknesses

Used alone, it is too late.

A challenge may arise after the analytical path has been lost, the model changed, scope shifted or relevant intermediate choices were not recorded.

The architecture depends on preconditioning and contemporaneous custody, not retrospective reconstruction. The Meta-Primer explicitly prohibits inferred or post-hoc intentionality.

The scorecard also cannot fairly assess a challenged answer unless the evidence required for scoring was captured when the answer was produced.

Verdict

Useful as an escalation mechanism, but not as the point at which scoring begins.

Challenge should trigger deeper review, not first-time reconstruction.

Recommended architecture

I would use four linked levels.

Level 1 — Matter Assurance Baseline

Created at intake and updated only when the matter configuration materially changes.

It covers:

- corpus and custody environment;
- scope and privilege rules;

- admitted sources and provenance baselines;
- server, model and primer versions;
- matter threshold;
- logging and integrity controls.

This provides an inherited baseline for later questions.

It is not a declaration that every future answer is defensible.

Level 2 — Evidence Item Status

Every accessed item receives lightweight, mostly automated status information:

Evidence ID

Source and custodian

Path and Message-ID

Content hash

Provenance state

Transformation history

Privilege/scope state

Duplicate/version relationships

Exception flags

No full Y or Z score is required unless the item itself becomes a disputed analytical object.

Level 3 — Question or Proposition Scorecard

This is the primary scorecard.

It is applied to every answer intended to be:

- relied upon;
- presented to counsel;

- used to form cross-examination;
- included in a report;
- disclosed to another party;
- or treated as a substantive finding.

The unit is not necessarily every conversational prompt. It is the **filed analytical proposition**.

For the court scenario, separate propositions might be:

1. Person A and Person B exchanged 28 unique messages.
2. Those messages formed eight distinct threads.
3. Five threads concerned Transaction Alpha.
4. The principal subjects included valuation and document timing.
5. One additional communication may have occurred outside the collected email set.

Each proposition can inherit the same run evidence but have different confidence and potentially different Y scores.

That is important: the count may be certain, while the “essence” may remain interpretive.

Level 4 — Challenge or Exception Review

Triggered when:

- opposing counsel challenges an answer;
- an evidence item is disputed;
- an alias is questioned;
- a provenance divergence appears;
- a trace-versus-log discrepancy arises;
- an axis falls below threshold;
- a bypass is requested;

- or the answer is to be elevated from internal use to court use.

This is not a new score invented after the event. It is a **deeper examination of the existing filed score and record.**

It may produce:

- confirmed score;
 - adjusted score;
 - qualification;
 - superseding analysis;
 - human release;
 - or rejection for the intended use.
-

The balance between cost and assurance

A useful principle would be:

Capture everything needed to score, but perform full scoring only at the level where reliance occurs.

That means:

- matter controls are continuously active;
- evidence-item status is mostly automatic;
- question-level scoring occurs when an answer is filed for reliance;
- challenge-level review occurs when dissonance or adversarial scrutiny appears.

This avoids both extremes:

- scoring half a million Enron emails individually;
- and waiting until cross-examination to discover that no reconstructible record exists.

Practical decision table

Level	What is assessed	When	Full X–Y–Z?
Matter	Environment and corpus controls	Intake and material change	Baseline X/Z, not answer-level Y
Item	Provenance and evidence status	On access or admission	No
Question/proposition	Analytical answer	Before reliance or filing	Yes
Challenge	Disputed item or answer	On exception or challenge	Reassessment and escalation

My recommendation

The core architecture should be:

Matter baseline + automated evidence passports + full scorecard for each filed question or proposition + deeper review when challenged.

The **question-level scorecard is the centre of gravity** because that is where evidence becomes reasoning and reasoning becomes something a human may rely upon.

The “perfect until it is not” principle then determines the presentation:

- clean inherited conditions remain quiet;
- the question-level result shows the score;
- any item-level or architectural exception becomes prominent;
- challenge activates deeper interrogation rather than retrospective invention.